



Compliance **AI**

Master in Management

2025–2026

Master's Thesis “Compliance.AI: A Business
Plan for EU AI Act Compliance in SMEs”

Francisco Cordeiro Batista

Supervisor: Professor Samira Reis

Madrid, June 2026

ABSTRACT

Compliance.AI is a business-to-business (B2B) software-as-a-service (SaaS) platform designed to help European small and medium-sized enterprises (SMEs) comply with the EU AI Act (Regulation EU 2024/1689). On 7 May 2026 the Council of the European Union and the European Parliament reached a provisional agreement under the Digital Omnibus to defer stand-alone Annex III high-risk obligations to 2 December 2027 (Council of the European Union, 2026). This plan is built around that window. This business plan sets out the strategic positioning, market opportunity, competitive structure, product architecture, and financial projections for a specialist compliance solution targeting a market gap left open by enterprise-focused competitors. The analysis is grounded in strategic frameworks including Porter's Five Forces, the resource-based view, and organisational ecology, positioning Compliance.AI as a focused competitor in the SME compliance technology market. The document is both a strategic roadmap and an academic thesis on market entry and sustainable competitive advantage in emerging regulatory technology sectors.

Keywords: EU AI Act, regulatory technology, SaaS, SMEs, compliance automation, market entry, business strategy, competitive positioning



CONTENTS

EXECUTIVE SUMMARY	1
1. BUSINESS IDEA AND STRATEGIC POSITIONING	3
Vision and Mission	3
The EU AI Act: A Plain-English Overview	3
1.1 Business Concept and Value Proposition	5
1.2 Market Position and Competitive Differentiation	6
1.3 Target Customer Segment and Needs	7
1.4 Strategic Positioning Framework	9
1.4.1 Value Proposition Alignment	9
1.4.2 Cost Differentiation: A Cluster Model	9
1.4.3 Resource Partitioning	10
1.4.4 Accumulated Customer Knowledge	10
1.5 Go-to-Market and Revenue Model	11
2. EXTERNAL ANALYSIS.	12
2.1 PESTEL Analysis	12
2.1.1 Legal and regulatory	12
2.1.2 Economic	13
2.1.3 Technological and social	15
2.2 Porter's Five Forces	16
2.2.1 Threat of new entrants	16
2.2.2 Threat of substitutes	16



2.2.3	Bargaining power of customers	17
2.2.4	Bargaining power of suppliers	17
2.2.5	Competitive rivalry	17
2.3	Market opportunity synthesis	18
3.	INTERNAL ANALYSIS	19
3.1	VRIO analysis	19
3.1.1	Valuable	19
3.1.2	Rare	21
3.1.3	Costly to imitate	21
3.1.4	Organised	22
3.1.5	From temporary advantage to sustained moat	22
3.1.6	Internal weaknesses	23
3.2	Capability stack	24
3.3	SWOT and strategic position synthesis	25
4.	MARKETING PLAN	27
4.1	Relevant Market	27
4.2	Pricing	29
4.3	Distribution Plan	30
4.4	Communication Strategy	31
5.	HR PLAN – THE TEAM	33
5.1	Workforce Planning and Job Design	33
5.1.1	Workforce planning	33
5.1.2	Job design	35
5.2	Managing Employee Competencies	36



5.2.1	Recruitment	36
5.2.2	Selection	36
5.2.3	Training	37
5.3	Managing Employee Attitudes and Behaviours	37
5.3.1	Performance	37
5.3.2	Compensation	38
5.3.3	Benefits	39
6.	OPERATIONS.	40
6.1	Service delivery and infrastructure	40
6.2	Client onboarding	41
6.3	Team operating model	41
6.4	Regulatory content pipeline	42
6.5	Compliance, partners and data governance	43
7.	FINANCIAL PLAN.	45
7.1	Investment and Funding Plan	45
7.2	Key Assumptions	46
7.3	Financial Statements	47
7.4	Unit Economics and Viability	48
7.5	Scenario Analysis and Contingency	49
8.	IMPLEMENTATION PLAN	50
8.1	The Critical Path: Build First, Sell Second	50
8.2	The Team Arrives in the Order the Work Demands	52
8.3	Security Compounds into a Sales Asset	52
8.4	The Sequence Funds Itself	53



9. CRITICAL RISKS AND CONTINGENCY PLAN	54
9.1 Risk Framework	54
9.2 Material Risks	54
9.2.1 Demand: enforcement delay and SME awareness (R1, R6)	54
9.2.2 Competition: generalist downmarket entry (R2)	56
9.2.3 Technology and liability: classification error (R3)	57
9.2.4 Key person: founder concentration (R4)	57
9.3 Sensitivity Analysis	58
9.4 Exit and Pivot Triggers	58
REFERENCES	59
APPENDIX	65
A.1 Risk Register	65
A.2 Key Assumptions Audit	68
A.3 Financial Model Assumptions	69
A.4 Five-Year Financial Statements	69
A.5 Scenario Detail	69



LIST OF FIGURES

1.1	Product architecture.	6
1.2	Strategic positioning map.	8
3.1	SWOT synthesis.	26
4.1	The Year 1 market funnel.	28
5.1	Year 1 organisational chart.	35
9.1	Risk heat map: ten risks scored by likelihood and impact.	55



LIST OF TABLES

1.1	Competitive comparison matrix.	7
3.1	VRIO summary on Day 1 (pre-launch).	20
4.1	Year 1 marketing KPIs.	32
5.1	Year 1 team compensation plan.	39
7.1	Seed funding structure.	46
7.2	Financial summary: base case (EUR).	47
8.1	Implementation roadmap.	51
9.1	Risk register: ten tracked risks.	66
9.2	Key assumptions, failure modes, and validation paths.	68
9.3	Core financial model assumptions.	70
9.4	Income statement: base case (EUR).	71
9.5	Cash flow statement: indirect method (EUR).	72
9.6	Balance sheet: year-end (EUR).	73
9.7	Four-scenario summary: Year 1 and Year 3 exit metrics (EUR).	74



ARTIFICIAL INTELLIGENCE USAGE DECLARATION

This declaration details the use of artificial intelligence (AI) tools during the preparation of this business plan, in accordance with UC3M academic integrity guidelines.

Tool Used

Claude (Anthropic), used as a research and analysis assistant.

Author's Declaration

I confirm that this business plan represents my own intellectual work. The AI tooling described above was used solely as a research and analytical aid. I take full academic responsibility for all content contained herein.

Francisco Cordeiro Batista

Master in Management, UC3M

Madrid, June 2026

AI Assistance: Scope of Use

AI-assisted tasks	Author's own work (no AI)
Literature and market research: identification, retrieval, and summarisation of academic sources and publicly available market intelligence. All cited sources were independently verified by the author before inclusion.	All written text, arguments, analyses, and conclusions appearing in this thesis. No AI-generated prose was submitted as original work.
Financial scenario modelling: generation of structured scenario inputs (base case, optimistic, and stress-test) for the financial projections in Chapter 7. The author defined all modelling assumptions, reviewed all outputs, and bears full responsibility for the projections as presented.	Strategic positioning, competitive framing, and all qualitative judgements throughout the document.
Go-to-market hypothesis testing frameworks: prototyping of A/B testing and pricing hypotheses used as analytical scaffolding. These outputs were not reproduced verbatim in the document.	Primary data collection, interpretation of regulatory sources, and all decisions regarding business model design and financial assumptions.
$\LaTeX$ coding and typesetting: structuring and debugging the document in $\LaTeX$, including table layouts, TikZ figures, and visual formatting of charts and diagrams. The author defined all content, structure, and design intent; $\LaTeX$ implementation was produced and iterated with AI assistance.	Selection and validation of all frameworks, models, and analytical content.



ACKNOWLEDGEMENTS

This thesis exists because of the people named below. Each one contributed something that could not have been replaced.

Professor Samira Reis supervised this project with rigour and patience in equal measure. She pushed me to ground every argument in evidence, to think in frameworks rather than opinions, and to treat the defence as a conversation, not a performance.

Miguel de Reina gave me access to a network I had no right to expect at this stage. His willingness to join the advisory board, share his experience building Google Spain and Portugal from the ground up, and challenge my assumptions about what enterprise sales in Southern Europe actually looks like made Chapter 4 and Chapter 8 sharper and more honest.

Professor Kurt Desender, even without being my supervisor, answered every question and every email I sent him. His help was fundamental.

Inês de Magalhães Monteiro never hesitated to help me with my Zotero bugs, APA 7th citation standards, and linguistic excellence.

This thesis is dedicated to everyone who believed the gap was real before there was any evidence it was.

Francisco Cordeiro Batista

Madrid, June 2026



EXECUTIVE SUMMARY

Compliance.AI will launch in Year 1 across Spain, Italy and Portugal as the firm will cover an unmet need in the European SME marketplace: affordable, self-service compliance with the EU AI Act (Regulation EU 2024/1689). No tool today combines exclusive Act focus, SME-level pricing and a self-service design. Not one.

EUR 193,000 consulting bill or EUR 199 per month. That is the trade.

The opportunity is a hard regulatory deadline with a structural gap behind it. Under the Digital Omnibus provisional agreement of 7 May 2026, stand-alone Annex III high-risk obligations take effect on 2 December 2027, backed by fines of up to EUR 15 million or 3% of global turnover (Council of the European Union, 2026; Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act), 2024). Meeting them through consultants costs EUR 50,000–250,000, 3 to 10% of a typical SME's revenue (Laurer et al., 2021). The enterprise governance vendors, Credo AI and Holistic AI among them, price at USD 30,000–150,000 and have no economic reason to move down. The SME tier is left open. Resource partitioning (Carroll, 1985) explains why they cannot follow Compliance.AI into it.

The platform delivers four modules mapped to the Act: a risk-classification engine (Article 6, Annex III), an Annex IV documentation generator, a conformity workflow (Annex VI), and post-market monitoring (Article 72). Pricing runs four tiers: a free classifier, EUR 199 Essentials and EUR 499 Professional self-service subscriptions, and a custom Enterprise tier at roughly EUR 1,499 per month. The Year 1 markets hold 434,000 SMEs in the 10–249 employee band, of which 7,000–15,000 fall in scope; an 8 to 10% reach factor yields 700–1,200 addressable firms (Eurostat, 2025).

The critical success factors for the firm in order to establish, sustain and appropriate a competitive advantage are accumulated customer knowledge, the advisor network, and post-market monitoring lock-in. Those three factors will determine the product's defensibility, increase customer retention, and make switching costs prohibitive for SMEs that



have built their compliance trail on the platform. The product offered is both scarce and relevant in the European SME market and will therefore help Compliance.AI to establish a competitive advantage.

The initial investment required to finance Compliance.AI is EUR 400,000, of which EUR 80,000 is founder equity. The firm will have a negative net income in the first two full years of operation; from Year 3 the net income is positive and increasing. The net income will be EUR (185,549), EUR (199,228), EUR 77,963, EUR 594,787 and EUR 1,435,519 respectively. The objectives of the firm are to be operationally profitable and free-cash-flow positive from Year 3 onwards. The required return on investment is 30% on the external angel tranche. EBITDA (earnings before interest, taxes, depreciation and amortisation) breakeven arrives at month 33 from inception, and the seed is recovered from cumulative cash flow during Year 5, with no second raise required.

In case of a pessimistic scenario, with a sustained 25% fall in monthly customer additions or a further delay beyond 2 December 2027, Compliance.AI should follow a focused exit strategy: reduce burn to a single founder on zero salary and pursue an acqui-hire or asset sale to a governance, risk and compliance (GRC) platform such as Vanta or Drata. In a more optimistic scenario, with 1.5 to 2 times base-case growth or an early France and Benelux launch, the firm should evaluate an accelerated Series A at a EUR 10–15 million pre-money valuation.

We built for this gap. We get there first.

1. BUSINESS IDEA AND STRATEGIC POSITIONING

Prove your AI is legal, without hiring a lawyer.

Vision and Mission

Vision. Every European SME that deploys AI should be able to prove it is legal, without hiring a lawyer or retaining a consultancy. Compliance.AI exists so that “can we use this AI system?” becomes a question software answers.

Mission. We turn the EU AI Act into code. Compliance.AI automates risk classification, documentation, and continuous monitoring so SMEs reach compliance in weeks, not months, at a fraction of the cost of professional services.

The EU AI Act: Overview

The EU AI Act is Europe’s first law written specifically for artificial intelligence (Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act), 2024). Its core idea is simple: the more harm an AI system could cause, the more the law demands of it. A spam filter and a system that decides who gets a mortgage both use AI, but only one can damage somebody’s life, so only one carries heavy duties. That single principle organises the entire law.

The Act sorts every AI system into four tiers. A small set of uses is banned outright, such as government social scoring and real-time facial recognition in public spaces; these are off the market, so no customer exists there. Most everyday AI counts as minimal risk and carries no obligation at all. Limited-risk systems carry one duty only: a chatbot, for

instance, must tell users they are talking to a machine. None of these three tiers creates work an SME would pay software to handle.

The business lives in the fourth tier. High-risk systems make decisions that materially affect people's lives, and they are the only tier that must pass a formal compliance process before reaching the market. Annex III lists the eight categories that qualify, including biometric identification, recruitment, credit scoring, and law enforcement. These are ordinary business tools, not exotic ones. A firm that screens CVs with software or scores loan applications with an algorithm is running high-risk AI, whether it realises it or not.

For those firms, three annexes define the actual work, and they run in sequence. Annex III answers the first question: is my system in scope? Annex IV defines what the firm must then produce: a technical file recording how the system works, what data trained it, and how its risks are managed. Annex VI explains how the firm proves all this: an internal-control procedure the firm runs itself, with no outside inspector. That last point carries the whole business. Because the firm certifies its own system, software can run the entire process, and almost every SME qualifies for this route.¹ The one exception is biometric identification without a harmonised standard, which still needs an external assessor and covers only a small fraction of SME users.

One date makes the obligation real. Under the Digital Omnibus provisional agreement of 7 May 2026 (Council of the European Union, 2026), high-risk obligations for stand-alone Annex III systems take effect on **2 December 2027**. From that day, a firm running an in-scope system without the required documentation faces enforcement.² Compliance.AI maps one module to each part of the work: classification to Annex III, documentation to Annex IV, the conformity workflow to Annex VI, and continuous monitoring to Article 72.

¹Article 43(1) requires notified body assessment only for biometric identification systems lacking harmonised standards (roughly 5% of SME users); Article 43(2) governs the self-certification route available to every other Annex III system. Article 72 creates the ongoing post-market monitoring duty: track model drift, log incidents, report against standards. Article 11(1) requires the Commission to publish a simplified documentation form for small and micro enterprises.

²The deferral from the original 2 August 2026 date is a provisional political agreement; it binds only once the Omnibus is formally adopted and published in the Official Journal. Until then the original date technically remains live (Marrache & Sieniewicz, 2026).

1.1 Business Concept and Value Proposition

Compliance.AI will launch as the firm will cover an unmet need in the EU SME marketplace: affordable, self-service compliance with the EU AI Act. No comparable tool combines exclusive Act focus, SME-level pricing, and a self-service design. Not one.

The market failure is structural. Between 2024 and 2025, AI adoption reached 17% among small firms and 30.36% among medium enterprises (Eurostat, 2025). Every adopter running an Annex III system takes on a legal duty, backed by fines of up to EUR 15 million or 3% of global turnover (Article 99(4)). Meeting that duty the traditional way costs EUR 193,000 upfront to build a Quality Management System, plus EUR 71,400 a year to maintain (Laurer et al., 2021). For a 30-person firm, that is an existential number.

EUR 193,000 consulting bill or EUR 199 per month. That is the trade.


Compliance.AI cuts the cost of compliance by nearly two orders of magnitude, turning a capital project into a monthly operating line. Four modules map to the Act's obligations (Figure 1.1):

- **Risk classification engine.** Determines whether an AI system is high-risk under Article 6 and Annex III. SMEs running recruitment screening or credit scoring often do not know their system is in scope. The classifier answers in minutes.
- **Documentation generator.** Produces the Annex IV technical file: system architecture, training data, risk management, and performance metrics, built around the simplified SME form that Article 11(1) mandates.
- **Conformity workflow.** Walks the firm through the internal compliance procedure step by step, generating the checklist, sign-off records, and declaration of conformity required before deployment.
- **Monitoring module.** Tracks the post-market duties under Article 72: model drift, incident reporting, and changes to harmonised standards.

The demand case rests on two findings. Seven in ten AI-considering SMEs cite lack of internal expertise as the main barrier to adoption (Eurostat, 2025). And a 2025 survey

Figure 1.1

Product architecture.



The four modules map one-to-one onto the Act’s obligations. The customer moves left to right; post-market findings and new authority guidance feed back into classification, the mechanism behind the accumulated-knowledge moat in Chapter 3.

of over 1,000 EU and UK technology SMEs found one-third had already removed or downgraded AI features under regulatory uncertainty, at average annual revenue losses of EUR 94,000–322,000 per firm (Association for Competitive Technology, 2025). Most SMEs do not pay the compliance bill. They strip the AI, or they ship it unprotected.

Compliance.AI replaces that choice with a subscription.

1.2 Market Position and Competitive Differentiation

The AI governance platform market reaches USD 492 million in 2026 and passes USD 1 billion by 2030 (Kornutick, 2026). Growth concentrates at the enterprise end, and the firms that dominate there have no reason to serve SMEs: their cost models prevent them from trying.

Credo AI is the largest incumbent, with USD 39.3 million raised (PitchBook, 2026) and a Forrester Wave Leader designation for AI Governance Platforms Q3 2025.³ Its pricing runs USD 30,000–150,000 per year, requires machine learning (ML) infrastructure integration, and assumes the buyer has a dedicated data science team (Pellerin, 2026). Holistic AI, a University College London spinout, targets the same enterprise segment at custom pricing (Holistic AI, 2025). Neither is a realistic option for a 30-person firm

³The Forrester Wave report sits behind a paywall; this classification is cited through Credo AI’s press release (Credo AI, 2025).

running a recruitment screen.

Two EU-native competitors sit closer. eyreACT (United Kingdom) targets SMBs with AI Act-specific risk scoring but remains in pilot with no published pricing.⁴ EuroComply covers five frameworks, including the AI Act, at roughly EUR 149 per month (EuroComply, 2026). Five regulations at that price means shallow coverage on each.

No incumbent combines exclusive AI Act focus, SME pricing, and self-service design. Compliance.AI does, at up to 60 times less than the enterprise platforms (Table 1.1).

Table 1.1

Competitive comparison matrix.

Criterion	Credo AI	Holistic AI	eyreACT	EuroComply	Big Four	Compliance.AI
EU AI Act depth	Moderate	Moderate	High	Moderate	High	High (exclusive)
Target segment	Enterprise	Enterprise	SME	SME	Enterprise	SME
Annual cost	\$30K–150K	Custom	Undisclosed	~€1,788	€50K+	€2,388–5,988
Self-service	Partial	Partial	Yes	Yes	No	Yes
Maturity	Established	Established	Pilot	Operational	Established	Pre-launch
Continuous monitoring	Yes	Yes	Partial	No	No	Yes
Conformity workflows	Yes	Partial	Yes	Partial	Yes	Yes

1.3 Target Customer Segment and Needs

Compliance.AI targets European SMEs with 10 to 249 employees (Commission Recommendation 2003/361/EC) that develop or deploy AI within the Act’s scope. A subset of them runs AI in regulated domains. **That subset is the market.**


Year 1 focuses on Spain, Italy, and Portugal. AI adoption there trails Northern Europe: Spain at approximately 20% and Italy at approximately 16%, against Denmark at 42.03% and Finland at 37.82% (Eurostat, 2025). That gap looks like a weakness. It is the entry strategy, for two reasons:

First, Spain hosts the EU’s first AI Act regulatory sandbox and created AESIA (Agencia Española de Supervisión de la Inteligencia Artificial), the pioneer national AI supervisory authority. Organisations that align early with local regulatory institutions gain legitimacy

⁴No public data was found for eyreACT’s pricing as of June 2026.

Figure 1.2

Strategic positioning map.



Bubble size indicates market maturity. Compliance.AI occupies the only position combining low customer cost with deep, exclusive EU AI Act focus. Incumbents cluster in the high-cost quadrant; the only competitor with comparable regulatory depth (eyreACT) remains in pilot.

and reduce uncertainty ahead of competitors (DiMaggio & Powell, 1983; Scott, 2001). Operating from Madrid, Compliance.AI calibrates its classification logic and templates to AESIA guidance as it publishes. A competitor headquartered in Dublin or Berlin cannot match that speed.

Second, competition follows adoption. High-adoption markets like Denmark and Finland draw dense enterprise-vendor activity; Southern Europe is thinner (Eurostat, 2025). We enter the markets the generalists have not reached, with a product built for them: native Spanish, Italian and Portuguese user interface (UI), AESIA-calibrated templates, and EUR pricing with no annual contract minimum.

SME awareness of the high-risk deadline is low, and that is the primary demand risk. Content marketing and AESIA sandbox proximity carry the awareness work in Year 1; the pricing tiers convert that awareness once it lands.

1.4 Strategic Positioning Framework

1.4.1 Value Proposition Alignment

Three customer pains define the market failure: the EUR 193,000 upfront cost (Laurer et al., 2021), the expertise gap that stops seven in ten AI-considering SMEs, and scope uncertainty, since many firms running Annex III systems do not know it. Each module attacks one of these pains directly.

Mapped to the Osterwalder et al. (2014) Value Proposition Canvas, time-to-compliance falls from months to weeks, cost drops to roughly one-tenth of consulting, and audit-ready evidence accumulates as the firm works. One assumption could weaken the fit: SMEs may not recognise their exposure before the deadline. The AESIA-aligned content engine in Chapter 4 is the answer to that.

1.4.2 Cost Differentiation: A Cluster Model

The cost advantage comes from clustering, not from the usual SaaS zero-marginal-cost fantasy.

Every SME using AI for recruitment screening faces identical Annex III obligations. Templates built for the first such customer serve the tenth at near-zero extra cost: the classification logic, the Annex IV fields, and the Annex VI workflow are the same. Credit scoring, customer-service automation, and document classification repeat the pattern. At scale across those clusters, cost of revenue settles at 18 to 40% of price, the EUR 500–800 per-account band the financial model carries (Appendix A.4). Not zero. Low enough to sustain EUR 199–499 monthly pricing, with gross margins that start near 60% under early human review and climb above 80% as the base scales.

Hybrid cost-differentiation strategies outperform Porter (1985) pure generic strategies in digital markets (Kim et al., 2004), a finding confirmed across 1,842 Spanish SMEs (Hernández-Perlín et al., 2024). Compliance.AI runs that hybrid in a segment the incumbents have declined to enter.

1.4.3 Resource Partitioning

Resource partitioning (Carroll, 1985; Carroll et al., 2002) explains why the enterprise incumbents cannot follow us downmarket. The theory says that when generalist firms compete for the centre of a market, they abandon its edges, and specialists thrive in the space left behind. That is exactly the structure here. Credo AI, Holistic AI, and the Big Four built their sales teams, onboarding, and cost bases around USD 30,000+ contracts. A EUR 2,388 deal, with acquisition payback that must clear in nine months, does not cover a single enterprise sales cycle. They are not absent from the SME tier by choice; their economics cannot price into it.

The pattern is familiar. HubSpot grew alongside Marketo, and Pipedrive against Salesforce, each holding a position the generalist had abandoned.

The window is 12 to 18 months, until a well-funded entrant recognises the segment. Speed to market and customer lock-in through annual contracts convert that window into retained share.

1.4.4 Accumulated Customer Knowledge

The durable moat is not the technology. It is the interpretive knowledge that compounds with every customer.

The first Spanish customer's Article 15 accuracy review teaches the platform how AESIA actually reads the clause. The second customer avoids six weeks of rework because the template already carries that reading. The tenth customer inherits nine customers' worth of regulatory interpretation, built into classification logic, documentation prompts, and workflow defaults.

Two properties make this hard to copy. It scales with customer volume, not engineering headcount. And a late entrant cannot buy it retrospectively: the interpretations live in audit outcomes, not in published guidance. Organisational learning of this kind accumulates through experience and embeds progressively over time (Argote, 2013).

The critical success factors for the firm in order to establish, sustain and appropriate a

competitive advantage are accumulated customer knowledge, the advisor network, and post-market monitoring lock-in. Those three factors will determine the product's defensibility, increase customer retention, and make switching costs prohibitive for SMEs that have built their compliance trail on the platform. The product offered is both scarce and relevant in the European SME market and will therefore help Compliance.AI to establish a competitive advantage. In the long run the firm can further sustain its first-mover advantage by building scale. To build scale the capabilities need to be durable, transferable and replicable (Grant, 2010).

1.5 Go-to-Market and Revenue Model

Compliance.AI enters through Spain, Italy, and Portugal in Year 1, converting AESIA sandbox proximity and EU-native content into early adopters. Pricing runs four tiers: a free risk-classification tool, the EUR 199 Essentials and EUR 499 Professional self-service subscriptions, and a custom Enterprise tier at roughly EUR 1,499 per month, all deliberately below the EUR 50,000 consulting minimum that prices out the mid-market. Distribution channels and the full acquisition architecture are detailed in Chapter 4. The value proposition of Compliance.AI consists of the following: automated risk classification and Annex III scoping, self-service Annex IV documentation generation, and continuous post-market monitoring under Article 72.

2. EXTERNAL ANALYSIS

The deadline is the demand.

This chapter examines the environment the firm enters. The PESTEL analysis (political, economic, social, technological, environmental and legal) explains the forces creating demand, the legal ones first because they make compliance spending mandatory rather than optional. Porter's Five Forces (Porter, 1980) then maps the competition those forces leave behind. Four claims close the chapter.

2.1 PESTEL Analysis

2.1.1 Legal and regulatory

The law turns compliance from a choice into an obligation. Once high-risk obligations apply, Article 99 of the EU AI Act exposes any firm operating high-risk AI to fines of up to EUR 15 million or 3% of global annual turnover, whichever is higher (Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act), 2024). The systems that count are ordinary ones: Annex III covers recruitment screening, credit scoring, biometric identification, education, and critical infrastructure. That is why the obligation reaches firms that never thought of themselves as AI companies. SMEs are not exempt. The statute caps their fine at the lower of the fixed amount and the turnover percentage, but the duty itself stays. A small firm running an automated CV screen is in scope on the same terms as a multinational.

The deadline has moved, but only provisionally, and the distinction matters. On 7 May 2026 the Council and the European Parliament agreed under the Digital Omnibus (COM(2025) 836) to defer Annex III high-risk obligations to a fixed later deadline (Marrache & Sieniewicz,

2026). The agreement binds only once formally adopted and published, so firms cannot yet treat the delay as settled. For the business, a delay changes the shape of demand, not its size. Instead of a panic rush against the original date, the same demand spreads across a longer, more orderly buying window. For a subscription product, that is the better outcome: steady adoption suits recurring revenue far better than one feast followed by famine.

Two missing pieces explain why the answer must be software rather than a checklist. The Commission has published neither the harmonised technical standards that tell firms how to comply (CEN-CENELEC JTC 21 now targets Q4 2026) nor the simplified Article 11(1) documentation form for SMEs. Without the standards, a firm cannot know what evidence will satisfy an auditor. Without the form, it has no template to fill in. A static PDF of guidance cannot track requirements that are still moving. Software that maps each obligation to the draft standards as they emerge, and updates the firm's documentation as the picture firms up, can. Only a maintained product tracks a moving target.

Enforcement is national, and the Member States move at very different speeds. Spain leads: AESIA published 16 compliance guides in December 2025 and runs the EU's first AI regulatory sandbox, with 12 pilot projects (Cooper & Oberschelp de Meneses, 2025b). Italy designated AgID (Agenzia per l'Italia Digitale) and ACN (Agenzia per la Cybersicurezza Nazionale) under Law 132/2025 in October 2025. Portugal designated ANACOM (Autoridade Nacional de Comunicações) in September 2025. France is eight months past its own deadline with no designated authority. The practical result is that the same Annex III clause may be read differently depending on which regulator reviews it. For the SME, that fragmentation is a real problem. For Compliance.AI, it is an advantage. A product built from Spain outwards, tracking each authority's published guidance, carries the local knowledge that matters most when an audit happens. An enterprise platform engineered for US multinationals cannot adapt to that level of national detail.

2.1.2 Economic

AI adoption across European SMEs is rising fast, and compliance capacity is not keeping up. Between 2024 and 2025, the share of small firms using AI rose from 11.21% to

17%, and among medium firms from 20.97% to 30.36%. The EU-27 average climbed from 13.48% to 19.95% in a single survey cycle (Eurostat, 2025). This matters because adoption and obligation are linked. A firm takes on a legal duty the moment it puts a regulated AI system into use, so each new adopter is a potential customer. The adoption curve is the demand curve for this product, and it is steepening.

Not every adopter falls in scope. The duty attaches only to AI that maps to Annex III, so the question is whether the high-adoption sectors and the regulated sectors overlap. They do. The fastest adopters are the sectors the statute targets: Information and Communication at 62.5%, Professional and Scientific activities at 40.4%, and Manufacturing at 17.3% (Eurostat, 2025). Applied to the Spanish, Italian, and Portuguese base of 10–249-employee firms, current adoption rates imply tens of thousands of AI-using SMEs across the Year 1 markets. The market is not a forecast. It already exists, and it is being measured.

The cost of the alternative makes the software case. The 2021 CEPS study priced a compliant Quality Management System at EUR 193,000–330,000 upfront, with the SME-specific band at EUR 144,350–247,150 and a further EUR 71,400 per year in maintenance (Laurer et al., 2021). No primary estimate has superseded it. Consulting is no cheaper: full engagements start at EUR 50,000 and reach EUR 150,000–250,000 (Laurer et al., 2021). For a typical SME with EUR 5 million in revenue, compliance alone would consume 3 to 10% of turnover. No firm of that size spends a tenth of its revenue on a single regulation. The demand is real; the existing supply is priced for a different customer.

The subscription closes that gap. The Essentials tier costs EUR 2,388 per year (Chapter 1). Against a EUR 50,000 consulting minimum, the platform pays for itself before February of its first year. The arithmetic is not marginal. It is the difference between a cost an SME cannot absorb and one it can.

The macro data confirm the trajectory. Compliance.AI is capitalising on an increasing trend of SME AI adoption intersecting with mandatory governance requirements. Regulatory technology (RegTech) is growing at a 21% compound annual growth rate (CAGR), from USD 24.34 billion in 2025 to a projected USD 112 billion by 2033 (Grand View Research, 2025); within it, the AI governance segment tracks towards USD 1 billion by 2030 (Kornutick, 2026). Capital is moving on the same view: Vanta raised USD 150 million

in July 2025 at a USD 4.15 billion valuation (Vanta, 2025). Investors are betting that the obligation, the adoption curve, and the pricing gap converge. This business plan makes the same bet.

2.1.3 Technological and social

The customer's problem is a skills gap, not a motivation gap, and that distinction shapes the go-to-market. Among EU enterprises that considered AI adoption and then declined in 2025, 70.89% cited a lack of relevant expertise as the primary barrier, ahead of legal clarity at 52.52% and data protection at 48.83% (Eurostat, 2025). SMEs are not refusing to comply; they are stuck. A firm with no compliance officer and no budget for one cannot read 113 articles of regulation and turn them into action. Self-service software is the only structure that fits both their budget and their headcount, because it supplies the missing expertise without requiring a hire.

An awareness gap compounds the skills gap. No pan-EU survey measures deadline awareness directly, but a September 2025 Bitkom survey of 604 German companies shows the scale of the problem: only 24% expect to be affected by the AI Act, while 32% believe they are not affected, 30% are still assessing, and 11% have not addressed the topic at all (Bitkom e.V., 2025). Awareness in the three Year 1 markets is likely lower still. The implication is direct: a firm that does not know the deadline applies to it will not respond to a sales pitch. The first job is education, not selling, a point developed in Chapter 4.

The cost of doing nothing is already on firms' income statements. A September 2025 survey of more than 1,000 technology micro, small and medium-sized enterprises (MSMEs) by the App Association and TechnoMetrica found over a third of EU and UK developers had stripped or downgraded AI features to cope with regulation, at losses of EUR 94,000–322,000 per year per affected firm, and EUR 160,000–453,000 for directly affected small-tech firms (Association for Competitive Technology, 2025). These firms are the most reachable cohort. They are paying the price of the regulation without receiving its protection, and a product that lets them comply and keep their features speaks to a cost they already feel.

One source of relief exists, and it is narrow. Article 57 sandboxes let firms test systems

under regulatory supervision, but Spain's ran only 12 projects in 2025. The mainstream SME, the firm with a deadline and no specialist to call, sits entirely outside that scheme and still needs a self-service answer.

2.2 Porter's Five Forces

Demand established, the next question is competition. The structural starting point comes from Chapter 1, §1.4: the enterprise incumbents cannot follow Compliance.AI downmarket, because their cost structures and procurement cycles are built for buyers SMEs do not resemble. The five forces show what that separation leaves on the field.

2.2.1 Threat of new entrants

The threat of new entrants is **moderate**, because the real barrier is not the obvious one. Capital is not the obstacle: a minimum viable product (MVP) ships for under EUR 500,000, within reach of any funded team. The genuine barrier is regulatory depth. Translating 113 articles and the draft CEN-CENELEC standards into deterministic software logic takes 12 to 18 months of focused work, and money cannot buy back that elapsed time. That figure measures a rival starting from zero. The firm's own six-month pre-launch sprint builds only the MVP core; the remaining depth accumulates through the launch year, which is why the barrier outlasts the build. First movers then compound the barrier through experience: every conformity assessment processed in Year 1 produces ground-truth data on how AESIA, ACN, and ANACOM actually interpret the statute's ambiguous clauses. A later entrant can match the funding but not the time. The window is open now and closes inside 18 months.

2.2.2 Threat of substitutes

The threat of substitutes is **low**, because each alternative fails the SME on a specific point. In-house compliance requires the expertise most SMEs lack. Consulting at EUR 50,000 and above does the job but prices out the mid-market. The Article 11(1) template, once

issued, will be a static document with no update mechanism and no monitoring, so it cannot track a regulation still in motion. Sandboxes serve only a handful of projects. None of these is self-service software at an SME price, which is the combination the market lacks.

2.2.3 Bargaining power of customers

The numbers reflect this. General SMB SaaS churns at roughly 40% a year (High Alpha, 2024; UserMotion, 2024); here, annual logo churn opens at the ~35% implied by the 3.5% monthly base rate and falls as switching costs accumulate. An SME that has logged nine months of Article 12 events cannot reconstruct that audit trail with a new vendor. The strategic task: win price-sensitive customers cheaply, then earn the embedded position that makes them stay.

2.2.4 Bargaining power of suppliers

Supplier power is **low and falling**, which protects margins. The two key inputs, cloud infrastructure and large language models (LLMs), are both getting cheaper. Cloud is commoditised. LLM pricing collapsed through 2025 and 2026: Mistral Small 3.1 runs at USD 0.10 input and USD 0.30 output per million tokens, and Claude Haiku 4.5 at USD 1.00 and USD 5.00 (AI, 2026). Mistral also offers EU-hosted inference, which solves the data-residency requirement for a product handling regulated information. At these input costs, technology never constrains pricing: cost of revenue is dominated by the early human-review layer, which is why gross margin starts near 60% in Year 1 and reaches 82% by Year 5 as that review automates.

2.2.5 Competitive rivalry

Rivalry **splits in two**, and the picture is opposite at each end. The enterprise tier is genuinely contested: Credo AI (USD 39.3 million raised), Holistic AI, Saidot, and IBM watsonx.governance all compete for Global 2000 accounts at contract values of USD 30,000–150,000 and above. None of them sells to SMEs, and none has an economic reason to

move down. The SME tier is close to empty. eyreACT is in pilot with no published pricing. EuroComply.app is live but without confirmed traction. Delbion is a consultancy, not a SaaS product. Vanta and Drata offer EU AI Act modules inside broader GRC bundles, but neither is priced or localised for Southern European SMEs. The contested fight is happening in a part of the market Compliance.AI does not enter, and the part it does enter has no established competitor. This is a vacancy, not a fight.

2.3 Market opportunity synthesis

Four claims summarise the external case, each resting on the evidence above.

Demand is mandatory. Article 99 fines, no published SME template, and no finalised harmonised standards mean SMEs cannot wait the problem out. A deferred deadline delays the obligation; it does not remove it.

Economics dictate the pricing model. Consulting at EUR 50,000–250,000 sits far beyond a typical SME’s compliance budget; EUR 199–499 per month sits inside it. RegTech growth at 21% CAGR (Grand View Research, 2025) and the AI governance segment’s path to USD 1 billion (Kornutick, 2026) move in the same direction. Price is the lever that converts a legal obligation into a purchase an SME can make.

The window is 12 to 18 months. The enterprise incumbents do not serve this segment and cannot move down into it. The SME-focused entrants are early, small, and mostly absent from Southern Europe. A specialist entering Spain, Italy, and Portugal now, with France and Benelux in Year 2, establishes its position before the segment crowds.

The constraint is execution, not market size. The market is large and growing, but deadline awareness runs low across all three Year 1 markets and distribution is fragmented. The binding constraint is reaching customers, not finding them. Content marketing and AESIA sandbox partnerships carry the acquisition load; Chapter 4 sets out how.

3. INTERNAL ANALYSIS

The moat is what we learn.

Chapter 2 showed that the SME tier of the EU AI Act compliance market has no funded specialist. This chapter asks a harder question: **does Compliance.AI have the resources to claim that seat and hold it?** The honest answer separates two kinds of advantage. Some resources give the firm a head start a competitor could eventually copy. One builds into a defence that strengthens the longer the firm operates. The analysis audits three resources against the VRIO filters (valuable, rare, costly to imitate, organised), then traces how the head start converts into a moat. The weaknesses are named in full, because a resource audit that hides its gaps is not an audit.

3.1 VRIO analysis

VRIO tests whether a resource can produce a sustained advantage. The resource must clear four hurdles in turn: valuable (it improves the firm's position), rare (competitors do not already have it), costly to imitate (rivals cannot easily copy it), and organised (the firm is set up to capture the value). Clear all four and the advantage lasts. Clear only the first two and the advantage is temporary. Table 3.1 maps the firm's three Day 1 resources against the four filters; the subsections explain each verdict.

3.1.1 Valuable

The first test is whether each resource improves the firm's competitive position. All three pass. The simplest way to see why is to ask what each one saves: the first saves the customer money, the other two save the firm money.

The customer value justifies the price. Vanta and Drata already proved that SMEs will

Table 3.1

VRIO summary on Day 1 (pre-launch).

Resource or capability	V	R	I	O	Position
Founder human capital (LLM/SaaS technical capability; AI Act product literacy)	Yes	Yes	Partial	Yes	Temporary advantage
Strategic relational capital (board advisor: first Google Iberia employee; 20-year Southern European enterprise network; warm AI Act-scoped client pipeline)	Yes	Yes	Yes	Partial	Temporary advantage
Accumulated customer knowledge (<i>built across 18–24 months</i>)	Yes	Yes	Yes	Partial	Sustained advantage in formation

“V” = valuable, “R” = rare, “I” = costly to imitate, “O” = organised to capture value. Rows 1 and 2 are temporary advantages protecting the acquisition window; Row 3 is a dynamic capability built across 18–24 months of customer cycles.

pay for compliance software, at average contracts of USD 13,500–18,000 (Sacra, 2024a, 2024b). The same logic applies here. The consulting baseline runs EUR 50,000–250,000 (Laurer et al., 2021); Compliance.AI replaces a EUR 193,000 fee with a EUR 2,388 annual subscription (Chapter 1). That trade is what the customer buys.

The founder profile is a cost advantage. A Data Science degree from NOVA IMS plus a Master in Management (UC3M, in progress) puts LLM integration skill and AI Act product literacy in one person. Most RegTech firms split those functions across two hires, and a compliance lead alone costs EUR 80,000–120,000 per year (Glassdoor, 2024). Removing that salary from the early income statement is what makes a EUR 2,388–5,988 annual contract viable from Day 1. A competitor paying two salaries cannot match the price without losing money on each account.

The board advisor is a pipeline that exists before any selling starts. He runs a Madrid marketplace consultancy whose clients operate automated pricing and advertising pipelines on Amazon, AliExpress and Mercado Libre, systems that fall within the Act’s scope. Those clients already trust the advisor, so the hardest part of selling compliance software,

getting the meeting, is solved before the first cold email is sent.

3.1.2 Rare

The second test is whether competitors already hold these resources. Here the distinction matters: a technical founder who can build LLM-powered compliance software is not rare on its own; several exist. What is rare is the pairing. No other SME compliance firm pairs a product founder with the first employee of Google Iberia, a person who built the Southern European enterprise sales network from zero over two decades. That pairing cannot be assembled by posting a job. And the rarity is not abstract: it is a concrete list of reachable buyers that the firm has and others do not.

3.1.3 Costly to imitate

The third test is the most demanding. A valuable, rare resource gives only a fleeting edge if a rival can copy it. Here the three resources split apart, and the split is the point of this chapter.

The technology is not costly to imitate. Frontier LLMs are commodity inputs, so a well-funded entrant can rebuild the classification engine, the documentation generator, and the conformity workflow within about 18 months (Sacra, 2024a). The software is a head start, not a wall.

The advisor's network is different. Relationships built over twenty years cannot be reproduced on a hiring timeline (Dyer & Singh, 1998). A competitor can recruit an advisor of its own, but not *this* network, because the value lives in the specific relationships rather than in the job title. Copying the role is easy. Copying the trust is not possible by hiring.

The third resource cannot be bought at all. Once the platform has processed 500–1,000 classification cases and observed how regulators responded to each, that knowledge sits in Compliance.AI's own logs. It appears in no public guidance, and engineering talent cannot reconstruct it, because it comes from having handled the cases, not from cleverness. A late entrant cannot read its way to this knowledge. It can only generate its own, slowly, after launching. The lag is built into the asset itself.

3.1.4 Organised

The final test is whether the firm is structured to capture the value its resources create. A resource the company cannot operationalise produces no advantage.

For the founder and advisor, the organisation is in place from Day 1. The founder runs product and early sales, the advisor opens doors across Southern Europe, and a Head of AI Act Compliance hired in Year 1 converts regulatory requirements into product logic and validates the classification outputs. Vanta and Drata succeeded with the same structure: an engineering-led core, domain advisors around it, and deeper regulatory expertise added once margins could support the cost (Sacra, 2024a, 2024b).

Customer knowledge scores only Partial, and the reason is timing. The case logging and feedback system ships from Day 1, so the raw data accumulates immediately. What is missing at launch is the analyst who turns 1,000 logged cases into a curated, queryable knowledge base. That hire comes later. The resource is captured in raw form from the start; the organisation to refine it arrives with scale.

3.1.5 From temporary advantage to sustained moat

The VRIO table resolves into one strategic story, and the whole business rests on it. Rows 1 and 2, the founder and the advisor, are temporary. Both decay towards parity within 18–24 months, because a competitor can in time hire a technical founder and assemble its own advisory board. Their job is not to defend the firm forever. Their job is to win the acquisition window.

What builds during that window does not decay. Every customer interaction that ends in a filed document and an observed regulatory response teaches the platform what each authority will and will not accept. This is ordinary learning-curve behaviour: error rates fall by 10–25% with each doubling of cases processed (Argote, 2013). For a late entrant, the consequence is decisive. After 18–24 months, the gap is made of accumulated case experience, and better engineers cannot manufacture it. The late entrant can wait or buy. Both options confirm the moat.

This logic feeds the firm's competitive-advantage thesis directly. The critical success

factors for the firm in order to establish, sustain and appropriate a competitive advantage are accumulated customer knowledge, the advisor network, and post-market monitoring lock-in. Those three factors will determine the product's defensibility, increase customer retention, and make switching costs prohibitive for SMEs that have built their compliance trail on the platform. The product offered is both scarce and relevant in the Southern European SME segment and will therefore help Compliance.AI to establish a competitive advantage.

3.1.6 Internal weaknesses

An honest audit names its weaknesses as plainly as its strengths. Four are real.

The first is no in-house regulatory expertise at launch. Until the Head of AI Act Compliance arrives in Year 1, every classification output needs human review before a customer can rely on it. That review compresses margins until case volume justifies the hire. The weakness is temporary, but real while it lasts.

The second is founder concentration. A single technical founder carries product, early sales, and customer success until the Year 2 hiring plan adds capacity, and that plan depends on winning enough Year 1 logos to fund it from revenue and the seed. The firm carries key-person risk through exactly the period when it can least absorb a setback.

The third is no customer references at launch. SMEs ask for peer references before signing a compliance contract, and a firm with no customers has none to give. The first ten customers pay a credibility tax, buying before social proof exists, so early acquisition is slower and more expensive than later acquisition.

The fourth is single-regulation exposure. In Year 1 the platform covers only the EU AI Act. NIS2, the Data Act, and the Cyber Resilience Act (CRA) sit on the roadmap but are not shipped, so the firm cannot cross-sell into its base until Year 2. A related risk sits inside this one: if the Commission publishes a free documentation template under Article 11(1), the documentation module loses its standalone value. The mitigation is bundle pricing. Classification, conformity workflows, and post-market monitoring hold their value even if templating is commoditised, so no single module has to survive for the business to.

3.2 Capability stack

VRIO explains why the resources are defensible. The capability stack explains how they become a working product. Four capabilities carry the value proposition, described in the order a customer meets them; the moat is what they add up to once the customer-knowledge layer starts compounding above them.

The first is the risk classification engine, which decides whether a customer's AI system is high-risk at all. A language model makes the initial pass and routes borderline cases to human review, because the model is fast but not yet trustworthy alone. The reason is concrete: LLM accuracy peaks at roughly 70% on legal reasoning benchmarks (Fan et al., 2026; Fei et al., 2024),⁵ high enough for clear cases, far too low to leave a compliance decision unsupervised. The human layer ships from Day 1, and the model's calibration improves with each case, so the share needing review falls over time.

The second is localised regulatory content, which makes the product work across different national regulators. It tracks each authority's published guidance and maps the differences into country-specific workflow templates: a Spanish customer and an Italian customer each see requirements matched to their own regulator. This is feasible at launch because Spain's regulator has published the most extensive national guidance in the EU (AESIA, 2025; Cooper & Oberschelp de Meneses, 2025a), giving the firm a concrete body of rules to build against from the start.

The third is the conformity workflow, whose design turns the standards delay into an asset. The first EU technical standard entered public enquiry in October 2025, eight months behind schedule, with publication expected in late 2026 (Soler2025; European Commission, 2025). Rather than wait, Compliance.AI ships interpretive workflows now and updates them as each standard is finalised. For a customer facing a deadline, an interpretive workflow that exists today beats a perfect one that does not. The delay is not a problem to manage. It is the window the firm operates in.

The fourth is post-market monitoring, the strongest lock-in in the stack, though its strength

⁵Both benchmarks test general legal reasoning rather than EU AI Act classification specifically; no Act-specific accuracy benchmark is yet available.




arrives late. On Day 1 it is a logging system, and a logging system is easy to replicate. Its value builds with time: by Month 12 it holds a performance and incident trail that a procurement team will not discard, because doing so means abandoning the firm's own compliance history. The capability that looks weakest at launch anchors the customer most firmly once data has accumulated.

3.3 SWOT and strategic position synthesis

The SWOT (strengths, weaknesses, opportunities and threats) pulls the internal and external pictures together so the position can be read at a glance. Figure 3.1 maps the Day 1 resources from this chapter against the environment from Chapter 2. The strengths shown are what exists today, not what the firm plans to build, which keeps the audit honest.

Figure 3.1

SWOT synthesis.



Strengths and weaknesses derive from the Day 1 resource audit; opportunities and threats derive from

Chapter 2.

The diagram reads as one argument, not four lists. The Day 1 strengths line up against the regulatory-deadline opportunity, and the standards delay has widened the window in which those strengths can be used. The race is to convert the temporary advantages into accumulated customer knowledge before they expire.

4. MARKETING PLAN

Own the search before the rush.

4.1 Relevant Market

The market case is a four-step funnel: how many SMEs exist, how many use AI, how many of those face the Act's heaviest obligations, and how many of those the firm can reach in Year 1. Each step narrows the count (Figure 4.1).

Step 1: the universe. Spain, Italy and Portugal hold roughly 434,000 companies with 10 to 249 employees (Eurostat, 2025). These firms are large enough to run software that matters and too small to keep lawyers on staff.


Step 2: the AI users. The AI Act only binds companies that use AI. At the 2025 EU adoption rate of roughly 20%, some 60,000 of those SMEs already do.

Step 3: the high-risk users. Within that group, the heavy obligations fall on companies whose AI appears in Annex III of the Act. Annex III is the law's list of high-risk uses: software that screens job applicants, scores creditworthiness, grades exams, or controls access to essential services. A company on that list must produce a technical file, pass a conformity assessment and monitor the system after launch. That paperwork is the problem the platform solves. Annex III concentrates in information and communication, professional and scientific services, and parts of manufacturing: 7,000 to 15,000 of the AI-using SMEs.

Step 4: the reachable buyers. Not every in-scope SME can be sold to in Year 1. Awareness of the deadline thins toward the smaller firms, budgets cluster at the larger ones, and France waits until Year 2. An 8 to 10% reach factor gives 700 to 1,200 SMEs the firm can realistically put an offer in front of: the serviceable addressable market, or SAM (Figure 4.1).

Figure 4.1

The Year 1 market funnel.



Each stage narrows the count from the full SME universe to the firms the company can realistically reach, with the country split beneath each stage. Author estimates from Eurostat (2025) Structural Business Statistics (SBS) and information and communication technology (ICT) usage data.

Why is this market still open? The established vendors cannot afford it. Credo AI and Holistic AI sell at USD 30,000–150,000 per contract per year through sales teams built for large enterprises; that cost of selling cannot make money on a EUR 199 customer, so the SME segment goes unserved (Chapter 1, §1.4). The high-risk enforcement deadline then gives every in-scope SME a date by which it must act. A gap becomes a queue of buyers, and no competitor is positioned to serve it.

The Year 1 target is deliberately small against that queue: 30 customers signed during the year, of which 24 remain active at year end once cancellations are netted out. That is 2 to 4% of the SAM. At the blended average revenue per account (ARPU) of EUR 339 per month (the average across the three paid tiers, derived in Chapter 7, §7.2), those 24 accounts exit Year 1 generating revenue at a pace of EUR 98,552 per year. France and Benelux add roughly 600–1,150 reachable SMEs from Year 2; the resource partitioning logic in Chapter 1, §1.4, keeps Spain, Italy and Portugal first.

4.2 Pricing

EUR 193,000 consulting bill versus a EUR 199 monthly subscription. That is the trade for every SME with no in-house legal team. Two comparisons close the case: what the alternatives cost, and what the SME already spends.

What the alternatives cost. Vanta and Drata are the closest SME-priced compliance platforms. Their entry tiers run USD 7,500–15,000 per regulatory framework per year, plus implementation fees of USD 10,000–25,000 in the first year (ComplyJet, 2025). A new customer therefore pays EUR 17,000–36,000 before signing a single compliant document. Compliance.AI Professional costs EUR 5,988 per year and treats the AI Act as the core product rather than a paid add-on. That is 65 to 83% below the all-in entry price of the nearest alternative.

What the SME already spends. Half of small businesses across Spain, the UK, France and Ireland spend EUR 1,000–50,000 a year on General Data Protection Regulation (GDPR) compliance (GDPR.eu, 2019), and the upper SME tier spends EUR 20,000–45,000 (Secure Privacy, 2026). Compliance.AI Professional costs 13 to 30% of that upper-tier figure. The buyer does not need to justify a new spending category; the subscription fits inside a budget line that already exists.

Two price anchors frame the rest of the space. Above, Credo AI and Holistic AI at USD 30,000–150,000 per contract do not compete for SME business. Alongside, EuroComply sells breadth across many regulations at roughly EUR 149 per month (EuroComply, 2026). The EUR 50 premium over EuroComply buys depth: the Annex IV technical file, the document pack the Act demands for every high-risk system, at a level of detail a many-regulation platform cannot deliver on any single law.

The critical success factors for the firm in order to establish, sustain and appropriate a competitive advantage in pricing are transparent EUR denomination, AI Act-specific depth, and self-service access up to the EUR 499 tier. Those three factors let an SME go from discovering an obligation to buying the solution in a single web visit.

4.3 Distribution Plan

Four channels carry the product to market. Each one meets the buyer somewhere they already are.

Content-led SEO. The SME buyer starts with a question, not a vendor search. A compliance officer who has just heard about the deadline types “is my recruitment tool high-risk under the AI Act” long before typing a product name. Today, page one for those searches is informational sites, not products. The firm publishes the answers in Spanish, Italian and Portuguese, tied to each national regulator’s guidance, starting from AESIA’s December 2025 publication (AESIA, 2025) and parallel material from AgID and ANACOM. Whoever answers the question becomes the obvious next step. The model is proven: Vanta draws roughly 250,000 organic visits a month this way, traffic worth about USD 128,000 if bought as advertising (Foundation Marketing, 2024).

Law firm and accounting partnerships. Mid-tier privacy practices in Spain (Garrigues, Pinsent Masons), Italy (Portolano Cavallo, ICT Legal Consulting) and Portugal (VdA, PLMJ) already advise SMEs on GDPR, and those clients now ask AI Act questions. The firms cannot answer at scale: drafting Annex IV documentation by hand does not fit a billable-hour model. The platform gives them a product to resell instead of work to turn away. Osano’s referral programme sets the playbook: a 10% client discount, joint marketing, and recurring revenue for the referring firm (Osano, 2025). The law firm keeps the client relationship and earns a fee; the platform gains a warm introduction from an adviser the SME already trusts.

Trade association distribution. The European DIGITAL SME Alliance represents 45,000 digital SMEs across 30 national associations and ships a free AI Act Conformity Tool with EIT Digital (European DIGITAL SME Alliance, 2025). That free tool feeds the funnel rather than competing with it: it tells an SME whether it has a problem, then stops. It writes no technical file and runs no workflow. A firm that discovers an Annex III obligation through the tool arrives needing exactly what the platform sells.

Article 57 regulatory sandboxes. Sandboxes are supervised test programmes that national regulators must run for high-risk AI projects. AESIA selected 12 such projects

in April 2025 (AESIA, 2025). Participants are the most qualified leads in the market: they hold a high-risk system, a live relationship with the regulator, and a deadline. The sandbox tests their system but does not produce its compliance documentation, so the platform enters as the tooling underneath the project. Every member state must run at least one sandbox, so the channel repeats in Italy and Portugal from Year 2.

The cost of these channels is measured as customer acquisition cost (CAC): the sales and marketing money spent to win one customer. Year 1 non-salary CAC is EUR 400 (a EUR 12,000 budget against 30 customers), which the EUR 199 Essentials subscription repays in 3.4 months (Chapter 7, §7.4). From Year 2, with a dedicated Account Executive on payroll, fully loaded CAC rises to EUR 800–1,200 and the payback period to roughly seven months. Both sit inside the KeyBanc 2024 benchmark of under 12 months for SMB software (European Digital SME Alliance, 2025; KeyBanc Capital Markets, 2024). Every channel above pushes CAC toward the lower bound: three of the four cost content and referral fees, not salaries.

4.4 Communication Strategy

Positioning. Compliance.AI is the EU-native, SME-priced AI Act compliance platform for Spanish, Italian and Portuguese businesses subject to Annex III. Figure 1.2 places it in the low cost-to-customer, high regulatory depth quadrant: the position incumbents cannot reach without dismantling their own sales model.

Differentiation. Five properties separate the firm from competitors that name SMEs in their marketing and then price them out:

- Native Spanish, Italian and Portuguese UI, documentation templates and support.
- EU-only data residency, with no Standard Contractual Clauses dependency on US infrastructure.
- An Article 11(1) simplified Annex IV template generator, the first SME-specific tool in the category.

- Templates aligned to AESIA, AgID and ANACOM guidance, updated within days of any authority change.
- Transparent EUR pricing on the website. No sales call up to the EUR 499 tier.

Customer journey. Regulator-aligned content and sandbox visibility create awareness. The free classification tool and deadline tracker turn a visitor into an active user. Annual contracts close the sale and cut churn 30–40% against monthly billing (OpenView Partners, 2022; ProductLed, 2024). The customer then renews for a simple reason: new authority guidance lands in the platform within days, not within a US vendor’s release cycle, and falling behind costs more than the subscription.

Year 1 key performance indicators (KPIs). Six metrics connect this chapter to the Chapter 7 financial plan (Table 4.1).

Table 4.1

Year 1 marketing KPIs.

Metric	Year 1 target
Freemium activations	600–1,000
Freemium-to-paid conversion	3–5%
Paid customers (eop, net of churn)	24 net (30 gross)
Blended ARPU	EUR 339/month
Non-salary CAC	EUR 400
Non-salary CAC payback	3.4 months

Targets aligned to the Chapter 7 base case and SMB SaaS benchmarks (Benchmarkit, 2025b; KeyBanc Capital Markets, 2024; OpenView Partners, 2022). Fully loaded CAC and payback from Year 2: see Chapter 7, §7.4. eop = end of period.

The funnel arithmetic is short. 600–1,000 free sign-ups at a 3–5% conversion rate produce 30 paying accounts. $30 \times \text{EUR } 339 = \text{EUR } 10,170$ in contracted monthly revenue. After 3.5% monthly churn, 24 accounts carry into the Year 1 exit at an EUR 98,552 annual recurring revenue (ARR) run rate. That is the target.

5. HR PLAN – THE TEAM

Three people. One deadline. Credibility on day one.

Three people close the gap that kills early B2B SaaS. A legal officer will not sign off on a compliance platform built by a data scientist with no AI Act credentials on the team. Italian buyers will not engage with a product that has never heard of AgID. And two markets cannot run on one founder. Year 1 hires one domain expert and one Italian commercial operator. That is the plan.

The model follows Lepak and Gowan (2017): workforce planning and job design, managing employee competencies, and managing employee attitudes and behaviours.

5.1 Workforce Planning and Job Design

5.1.1 Workforce planning

Year 1 runs on three people. The founder takes zero salary and covers engineering, product and co-sales across Iberia. Two paid hires fill the gaps the founder structurally cannot close: a **Head of AI Act Compliance** based in Lisbon or Madrid, and an **Italy Lead** based in Milan. A named advisory board sits alongside, anchored by Miguel de Reina, on phantom share agreements: contracts that pay a share of any future exit value without issuing real shares, so the advisors earn upside while the ownership register stays clean.

Three is not a compromise. Index Ventures (2017) benchmarks European seed-stage teams at roughly ten people on EUR 1 million raised; Compliance.AI runs three people on a EUR 400,000 seed. Vanta reached its first 600 customers on founder-led sales and an engineering-only team, adding governance specialists later (Sacra, 2024b). We add ours on day one. The difference is the rulebook. SOC 2 and GDPR are codified; the AI Act is

not. Waiting to hire credibility is the mistake Vanta could afford. We cannot.


Italy is the second constraint to remove. Law no. 132/2025, signed on 23 September 2025, is the first national AI statute inside the EU (Legge 23 settembre 2025, n. 132 – Disposizioni e deleghe al Governo in materia di intelligenza artificiale, 2025). It names AgID as the accreditation authority and authorises up to EUR 1 billion of public investment in Italian AI and cyber SMEs (Bonsignore et al., 2025). Credo AI is in Palo Alto; Holistic AI is in London. Neither has Italian-language documentation, an *avvocato* on staff, or a Milan office. A local-language specialist reaches a periphery the generalists cannot, and the Italy Lead is how that periphery converts in Year 1.

Advisory board. Miguel de Reina is the first and most active advisor. He was the first employee of Google Iberia, where he built the commercial operation from a single hire to a team of 50+ between 2003 and 2006. He now runs Brainy S.L., a Madrid marketplace consultancy whose clients run automated pricing and advertising pipelines on Amazon, AliExpress and Mercado Libre. Those pipelines fall within AI Act Article 4 and, depending on configuration, Annex III. His client base is therefore a warm prospect list before cold outreach starts. Miguel will also source the Italy Lead through his network, cutting search risk on the most commercially critical Year 1 hire. His terms: a phantom share agreement equivalent to 2.0% of exit value, vesting monthly over two years, with a 30-day termination clause. Phantom shares were chosen over a standard advisor share grant (FAST) to keep the cap table clean ahead of a future Series A (Index Ventures, 2017). Two to three further advisors with EU AI governance and data-protection backgrounds fill the remaining 1.5% of the 3.5% pool at 0.25–0.5% each. Figure 5.1 shows the resulting Year 1 structure: two reporting lines into the founder, with the advisory board alongside.

Year 2 adds up to four roles. Each hire fires on a measurable trigger, not a calendar date, and the base case lands three of them inside Year 2. An engineer joins when the founder's capacity blocks the product roadmap for more than two weeks. An Iberia Account Executive (AE) follows once the founder and Italy Lead run at full sales capacity. A Customer Success (CS) Manager joins once ARR crosses EUR 240,000 and onboarding outgrows founder bandwidth (Chapter 6, §6.2); the base case crosses that threshold during Year 2. The French-speaking commercial lead joins three months before France and Benelux open.

Figure 5.1

Year 1 organisational chart.



Year 2 hires shown as dashed ghost roles. Solid lines indicate Year 1 reporting relationships; long dashes indicate advisory and trigger-based Year 2 additions. Each Year 2 role fires on a measurable trigger, not a calendar date.

5.1.2 Job design

Founder / CEO. Engineering, product, and Iberia co-sales in Year 1. Territory: Portugal in full and one third of Spain. Sales runs in Spanish, Portuguese, and English. Technical scope covers all four Chapter 1 modules. One hybrid data-science / Master in Management (MiM) profile covers all this.

Head of AI Act Compliance. The hire that makes the Essentials and Professional tiers credible. They draft Annex IV templates for high-risk systems, own the regulatory-update queue against AESIA, AgID, and ANACOM guidance, and run the customer trust calls the founder cannot cover alone. Profile: four to nine years of data-protection or AI-governance practice, AI Governance Professional (AIGP) certified or willing to certify within ninety days, AI Act fluency shown through publications or advisory work, working Spanish and English. Lisbon is the preferred base; Portugal's 24.75% employer social-charge load, the tax an employer pays on top of salary, is the lowest across the three Year 1 markets (PwC Worldwide Tax Summaries, 2025b).

Italy Lead. A commercial hire, not a compliance one. The Italy Lead owns all Italian sales and marketing plus two thirds of Spain, working alongside the founder on Span-

ish mid-market outreach. The profile is a B2B SaaS operator with proven Italian enterprise relationships, sourced through the advisory network. The Head of Compliance and platform content supply the regulatory depth; the Italy Lead supplies market access and closes deals. No *avvocato* qualification required. The ability to sell to a CFO in Italian is. Milan-based, on a low fixed base plus commission on net new ARR, so the upside is tied to results.

5.2 Managing Employee Competencies

5.2.1 Recruitment

The Year 1 recruitment budget is zero in agency fees. LinkedIn Recruiter handles outbound; founder and advisor referrals carry inbound; the Italy Lead search is already warm through Miguel. Both roles open in parallel in the first quarter, with three to four months allowed to fill each.

Supply is thin. The AIGP-certified population is small globally (International Association of Privacy Professionals, 2025), and Iberia and Italy hold a fraction of it, with banks like Mediobanca and Intesa Sanpaolo bidding for the same Milan candidates. Compliance.AI cannot match those base salaries. The pitch is equity, category ownership, and the chance to build the market first.

5.2.2 Selection

Selection follows the person-organisation fit approach in Lepak and Gowan (2017). Strategic context comes first. A specialist firm needs candidates who prioritise regulatory aptitude and product instinct over law-firm pedigree.

The Head of Compliance screening is short and produces usable work. A 30-minute regulatory-fluency call filters out generalists. A two-hour timeboxed work sample asks the candidate to draft an Annex IV technical documentation template for a hypothetical SME running a CV-screening AI. The output is a content asset whether or not the candidate is hired.

The Italy Lead work sample is different. No documentation templates. The candidate receives a one-page product brief and runs a simulated discovery call with the founder playing a sceptical Italian operations manager. Regulatory accuracy is not the test. Commercial instinct is. Two reference calls with prior B2B peers close the process. No HR consultancy. No assessment platform.

5.2.3 Training

Year 1 training centres on two certifications. The Head of Compliance pursues the IAPP AI Governance Professional credential (EUR 1,500–2,000; International Association of Privacy Professionals, 2025); the founder pursues PECB ISO/IEC 42001 Lead Auditor (EUR 750–2,300). ISO/IEC 42001, the first AI Management System standard, is what procurement teams will ask for the way they now ask for SOC 2. Year 1 training spend: about EUR 4,000.

Article 4 of the AI Act, in force since 2 February 2025, requires providers and deployers to ensure their personnel hold demonstrable AI literacy. We run an internal AI-literacy programme as both regulatory obligation and sales proof point. The team selling SMEs an AI-literacy template runs one itself.

5.3 Managing Employee Attitudes and Behaviours

5.3.1 Performance

The team works in quarterly objectives with measurable key results (OKRs). Three top-level Year 1 OKRs frame the work: 24 net paid customers (30 gross), an ARR run rate of EUR 98,552 at the Year 1 exit, and zero customer regulatory escalations unresolved beyond five working days. The first two link to Chapter 4 Table 4.1 and the Chapter 7 base case. The third exists because reputation is the main sales asset in a 700-firm reachable market: one mishandled escalation costs three new customers through reference damage.

Role-specific KPIs sit underneath. Net new ARR for the founder, against the one-third Spain and full Portugal territory. Template velocity and a customer trust score (NPS) for



the Head of Compliance. Italian net new ARR and two-thirds Spain ARR co-ownership for the Italy Lead. The Year 1 net revenue retention target is 95%: the firm aims to keep 95 cents of every revenue euro from the prior year after churn and upgrades, in line with the Appendix A.3 assumption and the SMB SaaS benchmark range (Benchmarkit, 2025b; KeyBanc Capital Markets, 2024).

5.3.2 Compensation

The founder takes zero salary in Year 1. The policy reviews at the earlier of three events: ARR crossing EUR 240,000, any external financing event, or month 18. The trigger is fixed in advance. Index Ventures (2017) is direct on the point: extended unpaid founder time correlates with distressed bootstrapped exits, and equity to a serious early hire buys more runway than skipping a salary.

The two hires are paid to match their roles. The Head of Compliance sits at EUR 50,000 base, defensible against Glassdoor and ERI SalaryExpert ranges for Iberian data-protection officers in 2025–2026, with a 1.0% equity grant on four-year vesting and a one-year cliff. The Italy Lead sits at EUR 30,000 base plus a 10% commission on the first-year value of every new contract closed, paid quarterly in arrears, with the same 1.0% equity grant. At the base-case sales target the Italy Lead costs the firm about EUR 57,000 (EUR 46,450 fixed plus roughly EUR 10,550 commission) and earns about EUR 40,550. The structure rewards output, not presence.

Advisory equity follows Carta first-hire benchmarks (Carta, 2024). Miguel de Reina holds 2.0%; the remaining 1.5% pool splits across two to three domain advisors at 0.25–0.5% each, all on phantom share agreements with two-year monthly vesting and a 30-day termination clause. Real dilution at the end of Year 1 is therefore 2.0%, from the two employee grants. The 3.5% advisory pool is phantom and touches no shares on the cap table.

Table 5.1

Year 1 team compensation plan.

Role	Base (EUR)	Employer load	Other	Total fixed (EUR)	Equity
Founder / CEO	0	0	2,000	2,000	existing
Head of Compliance (Lisbon)	50,000	12,375 (24.75%)	5,500	67,875	1.00%
Italy Lead (Milan)	30,000	8,700 + 2,250 TFR	5,500	46,450	1.00%
Year 1 total (fixed)	80,000	23,325	13,000	116,325	2.00% real + 3.5% phantom (adv.)

Employer-load percentages: Portugal social security 23.75% plus accident insurance and wage guarantee fund bringing the all-in load to 24.75% (PwC Worldwide Tax Summaries, 2025b); Italy INPS plus accident contributions ~29%, with a separate 7.5% *Trattamento di Fine Rapporto* (TFR) severance accrual (PwC Worldwide Tax Summaries, 2025a). “Other” covers private health, equipment and certifications. Italy Lead commission (10% of closed first-year ACV, paid quarterly) is variable and excluded from the fixed total; on-target earnings (OTE) ceiling at base-case acquisition target is approximately EUR 57,000.

5.3.3 Benefits

The benefits stack matches early-stage European tech practice: private health insurance at EUR 600–1,200 per employee per year, a one-off EUR 2,000 equipment budget, and the EUR 1,500 annual learning budget that funds the certifications above. Remote-first by default. Statutory thirteenth and fourteenth-month payments in Iberia and the Italian TFR severance accrual sit inside Table 5.1.

The critical success factors for this team in order to establish, sustain and appropriate a competitive advantage in the SME compliance market are domain credibility, local-language commercial presence in Year 1 markets, and founder-led engineering velocity. Those three factors will determine the trust signal the firm projects, close the localisation gap in Italy, and keep the product moving while sales scales. The HR risks – founder concentration, zero-salary expiry, and specialist scarcity in Madrid and Milan – carry into Chapter 9.

6. OPERATIONS

Software ships. Trust is built.

Compliance.AI sells software, so most of operations is the software running itself. What remains human is small and decisive. Three questions decide whether the firm works in Year 1: how a new customer becomes productive in the first month, how a three-person remote team stays coordinated, and how a pre-revenue company earns enough trust to close its first thirty customers. This chapter answers each in turn.

6.1 Service delivery and infrastructure

The platform runs on Amazon Web Services (AWS) in Frankfurt. A second, ready-to-activate copy waits in Dublin in case Frankfurt fails. Both locations sit inside the EU, and that is the point. Customers store the evidence of their own compliance on the platform. Hosting that evidence outside the EU would create a new cross-border data problem on top of the one the customer pays to remove. EU hosting is therefore fixed in the architecture, not negotiated contract by contract.

Customers normally share the same infrastructure, which keeps costs low. A customer paying more than EUR 15,000 per year can have its data held on separate, isolated servers instead. Isolation is expensive, so only the largest accounts need it and pay for it.

The reliability promises are simple. The contract guarantees that the service is available 99.5% of the time; the firm runs to a stricter internal target of 99.9%, so the promise has headroom. If the system fails, two limits apply: at most one hour of data can be lost, and the service returns within four hours. Encrypted copies of all customer data are saved every day in both regions to keep those limits real.

The stakes behind these numbers are specific. The AI Act obliges companies to keep logs of how their AI systems behave (Article 12). Logs record the past, and the past cannot

be re-recorded. A customer with nine months of logs on the platform holds a record it cannot rebuild anywhere else. The backups protect that record, and the record keeps the customer: leaving Compliance.AI means leaving the audit trail behind.

6.2 Client onboarding

Onboarding is the first compliance deliverable, not an administrative step. The platform creates the account automatically; the founder builds the relationship.

The sequence is fixed. Day 0: the contract is signed, payment is set up through Stripe, and the account goes live with no waiting. Days 1–3: the founder runs a 60-minute call with a set agenda: classify the customer's AI systems by risk, import the full system inventory, and confirm which systems fall under the Act's high-risk list (Annex III). The customer leaves the call with a complete system profile and a named compliance owner inside its own organisation.

Week 2: the platform produces the first draft of the technical documentation the Act requires (the Annex IV file). The customer reviews flagged items inside the platform itself, with no email chains and no version confusion. Week 4: a 30-minute checkpoint call closes open items and any new requirements from AESIA, the Spanish supervisory authority. Month 3: the first post-market monitoring cycle (Article 72) completes and the customer holds a full audit trail. The renewal conversation starts there, nine months before the contract anniversary.

The founder runs every onboarding call for the first thirty customers. The Italy Lead takes Italian onboardings from the second quarter. A dedicated Customer Success hire takes over once annual recurring revenue passes EUR 240,000, which the base case places early in Year 3. Until then, the founder owns every activation.

6.3 Team operating model

The firm runs fully remote in Year 1. The entire office budget is EUR 150 per month: one co-working day per week at Utopicus Gran Vía in Madrid. Client meetings happen



at the client's premises or at a desk booked for the day. Credibility comes from a working product and a founder who shows up, not from a lease.

The weekly rhythm does not move. Monday: a 30-minute all-hands with a written record in Notion. Wednesday: a 60-minute regulatory review between the founder and the Head of Compliance. Friday: a 45-minute pipeline review between the founder and the Italy Lead. Client calls fill Tuesday and Thursday afternoons, capped at three per day, so the founder writes code in the mornings.

Miguel de Reina (Chapter 5) meets the founder once a month to review the pipeline, Southern European introductions, and positioning. His phantom-share agreement costs the firm no cash; a EUR 5,000 advisory stipend sits in the Year 1 payroll reconciliation (Appendix A.3).

Travel is deliberate and cheap. Any account worth more than EUR 7,000 per year gets one in-person founder visit at signature. The EUR 15,000 threshold from §6.1 is separate: it governs where data lives, not who gets a visit. Italy gets a founder trip to Milan each quarter; Portugal is handled remotely through the channels in Chapter 4. The full Year 1 travel budget is EUR 1,800.

6.4 Regulatory content pipeline

The platform updates on two clocks. Product code ships weekly. Regulatory content ships within 14 days of a new rule or guidance document being published. That second clock is a customer-facing promise, sold against US competitors whose content updates run quarterly.

The 14 days follow five fixed steps. Day 1: an analyst logs the new document. Day 3: a Spanish or Italian-qualified lawyer confirms whether it is legally binding and whom it covers. Day 7: the team drafts the rule changes and the customer checklist updates. Day 10: a second reviewer checks the work and tests it against real customer scenarios. Day 14: the update goes live with a changelog inside the platform.

The pipeline has material to process from day one. AESIA published 16 guidance documents in December 2025, and these form the launch baseline (AESIA, 2025). The

European harmonised standards now in development, led by prEN 18286 on quality management, enter the same pipeline as they publish (European Commission, 2025).

6.5 Compliance, partners and data governance

A compliance company must be visibly compliant itself. Two perimeters apply from day one.

The first is GDPR. Standard data processing agreements (Article 28) govern every customer relationship. The security measures GDPR requires (Article 32) are in place before the first invoice: stored data is encrypted, connections are encrypted, access is limited by role, and customers are notified of any incident within 72 hours. On top of GDPR, the firm pursues SOC 2, an independent security audit that larger buyers expect from any software vendor. Observation starts in Q2 of Year 1, the first report (Type I) arrives by Q4 of Year 1, and the stronger Type II report by Q3 of Year 2. The combined budget for the platform work, the audit, and penetration testing is EUR 60,000 to 90,000, spread across Years 1 and 2.

The second perimeter is the AI Act applied to Compliance.AI itself. The platform recommends risk classifications; a named human must approve each recommendation before it takes effect. That single design choice keeps the platform on the right side of the Act's line: it is a decision-support tool under Article 14, not an automated decision-maker that would itself count as high-risk (Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act), 2024). Without the human gate, Compliance.AI would carry the same monitoring obligations it sells to its customers. The choice is documented in the technical file and in the customer licence agreement.

Five vendors carry the stack. AWS Frankfurt is the sole cloud provider until revenue passes EUR 5 million. Stripe handles card and Single Euro Payments Area (SEPA) payments; GoCardless handles direct debit on annual contracts. Integrations with the accounting and enterprise resource planning (ERP) systems SMEs already use deploy in market order: Holded and Sage for Spain, TeamSystem for Italy, Primavera BSS for Portugal.



The infrastructure is rented. The trust is not.

7. FINANCIAL PLAN

Profitable by design, not by luck.

The initial investment required to finance Compliance.AI is EUR 400,000, of which EUR 80,000 is founder equity. The firm will have a negative net income in the first two full years of operation. From Year 3 the net income is positive and increasing. The net income will be EUR (185,549), EUR (199,228), EUR 77,963, EUR 594,787 and EUR 1,435,519 respectively.

The objectives of the firm are to be operationally profitable and free-cash-flow positive from Year 3 onwards. The required return on investment is 30% on the external angel tranche. The full five-year statements sit in Appendix A.4; this chapter explains, in plain terms, where the money comes from, where it goes, and why the numbers hold together.

7.1 Investment and Funding Plan

An SME that reaches the same compliance position through consultants pays around EUR 193,000. The platform charges EUR 199 per month (Laurer et al., 2021). That price gap is the demand the seed finances.

The logic of the raise is simple. The firm spends more than it earns for two years, so it needs enough cash to cover those losses plus a safety margin. Operations consume EUR 172,143 of cash in Year 1 and EUR 171,641 in Year 2, then generate cash from Year 3. Add two years of equipment and software purchases and the business needs EUR 370,000 before it funds itself. A EUR 30,000 cushion covers the end of Year 2, when the bank balance is thinnest. EUR 370,000 of need plus EUR 30,000 of cushion is the EUR 400,000 raised. Not a euro more.

The four sources behave differently. Founder equity and the angel tranche buy shares: together they put EUR 325,000 of paid-in capital on the balance sheet. The ENISA line is

Table 7.1*Seed funding structure.*

Source	Amount (EUR)	%
Founder equity	80,000	20%
External angel / pre-seed	245,000	61.25%
ENISA Digital Credit Line (Spain)	50,000	12.5%
Startup Portugal voucher programme	25,000	6.25%
Total seed	400,000	100%

a loan from the Spanish state at 4 to 6% interest, with no capital repayments until the end of the term (ENISA — Empresa Nacional de Innovación, 2025). The Startup Portugal voucher is a grant the firm never repays (Startup Portugal, 2025). The loan and the grant bring in EUR 75,000 without the founder selling a single extra share.

7.2 Key Assumptions

Four assumptions drive every number in this model. First, price. The average customer pays EUR 339 per month, a weighted average of the tier mix priced in Chapter 4: 70% of customers on the EUR 199 Essentials tier, 25% on Professional, 5% on Enterprise. Second, churn, the share of customers who cancel each month. It starts at 3.5% and falls to 1.8% by Year 5, the benchmark for software sold to small businesses (UserMotion, 2024). Third, gross margin, the part of each euro of revenue left after the direct cost of serving a customer. It starts at 60%, because a human still reviews early documentation, and reaches 82% by Year 5 as software takes over that work (Benchmarkit, 2025a; Lighter Capital, 2024). Fourth, volume: the base case adds 30 new customers in Year 1. The full assumption set, with sources and confidence ratings, sits in Appendix A.3.

One calendar note: Year 1 starts at the August 2027 platform launch (Chapter 8). The pre-launch period, July to December 2026, builds the product and earns nothing. Year 1 therefore contains only five selling months, August to December 2027.

The 30-customer base case is the honest number. The founder closes ten customers, the

Italy Lead converts ten inbound from Milan, and organic content brings the final ten. Churn takes its share: 30 new customers become 24 by year-end, worth EUR 49,276 of recognised revenue. The model assumes no enterprise deals and no advisory windfalls.

7.3 Financial Statements

Each statement answers a different question. The income statement asks whether the firm earns more than it spends. The cash flow statement asks whether there is money in the bank. The balance sheet checks that everything adds up. Together they tell one story: revenue climbs from EUR 49,276 to EUR 3.7 million across five years, profit arrives in Year 3, and cash never runs out on the way there. Table 7.2 carries the headline; the full statements sit in Appendix A.4.

Table 7.2

Financial summary: base case (EUR).

	Y1	Y2	Y3	Y4	Y5
Revenue	49,276	308,772	989,963	2,134,787	3,708,519
EBITDA	(181,049)	(189,228)	88,963	606,787	1,453,519
Net income/(loss)	(185,549)	(199,228)	77,963	594,787	1,435,519
Cash, end of year	219,857	30,216	125,885	735,475	2,177,571

Income statement. Revenue follows the customer count: 24 customers at the end of Year 1, 301 by Year 3, 843 by Year 5. The gross margin widens from 60% to a Year 3 peak of 84%, then settles at 82% as Enterprise accounts add bespoke review hours. Net income turns positive at EUR 77,963 in Year 3 and compounds from there. The firm also pays no tax until Year 4. Spanish law lets a company use past losses to cancel future taxable profit, and the EUR 384,777 lost in Years 1 and 2 shields the Year 3 profit and part of Year 4. From Year 4 Spain's reduced 15% rate for new companies applies (Price-waterhouseCoopers, 2025); the figures shown are pre-tax, with the tax effect quantified in the Appendix A.4 footnote and funded from operating cash flow.

Cash flow. The cash line looks odd at first sight. The firm loses EUR 185,549 in Year 1

yet ends the year with EUR 219,857 in the bank. Two things explain it. The EUR 400,000 seed lands that year. And subscription customers pay in advance, so cash arrives before the revenue is earned; accountants book the difference as deferred revenue. That prepaid cash grows with the customer base and funds part of the losses from inside the business. The bank balance bottoms at EUR 30,216 at the end of Year 2, the tightest moment of the plan, then climbs.

Balance sheet. Assets equal liabilities plus equity in every year, the proof that the three statements agree. Cash is the largest asset and deferred revenue the largest liability, the normal shape for a software firm that owns little physical equipment. Accumulated losses push equity slightly negative in Year 2, the same trough the cash statement marks, and equity recovers from Year 3 as profits are retained.

7.4 Unit Economics and Viability

Unit economics answer one question: is a single customer worth more than the cost of winning them? Two numbers decide it. Lifetime value (LTV) is the gross profit one customer generates before cancelling. Acquisition cost is what the firm spends to win them. At the EUR 199 entry tier, lifetime value is EUR 3,411, set by the price, the 60% gross margin and 3.5% monthly churn. The founder wins that customer for about EUR 400 in Year 1, so the customer is worth more than eight times what they cost, and the cost is recovered in 3.4 months of subscription payments. From Year 2 a salaried salesperson raises the fully loaded acquisition cost to EUR 800–1,200 and the recovery period to roughly seven months (Bessemer Venture Partners, 2024; Hays, 2025). Both years sit well inside the 3:1 lifetime-to-acquisition ratio investors treat as the minimum.

Two tests decide whether the plan funds itself. EBITDA, the operating profit before interest, tax and depreciation, breaks even in Year 3, 33 months from inception. The EUR 400,000 seed is recovered from cumulative cash flow during Year 5, by which point the operating margin reaches 39%. The firm reaches profitability on the seed alone. No second raise is required.

7.5 Scenario Analysis and Contingency

The base case turns on one number: customers won in Year 1. Four scenarios stress it, from a bull case of 60 customers to a downside of 8, with the full grid in Appendix A.5. The capital need moves with the outcome. The bull case breaks even in 18 months on EUR 350,000, the base case in 33 months on EUR 400,000, the bear case needs EUR 500,000 and the downside EUR 600,000.

In case of a pessimistic scenario, with a sustained 25% fall in monthly customer additions or a further delay to the high-risk enforcement deadline, Compliance.AI should follow a focused exit strategy: cut spending to a single founder on zero salary, end the Italy Lead engagement, and sell the firm or its technology to a larger compliance platform such as Vanta or Drata seeking EU AI Act content. In a more optimistic scenario, with 1.5 to 2 times base-case growth or an early launch into France and Benelux, the firm should evaluate an accelerated Series A at a EUR 10–15 million valuation before new money, consistent with market multiples for compliance software at EUR 1 million of annual recurring revenue (Bessemer Venture Partners, 2024).

The critical success factors for the firm in order to establish, sustain and appropriate a competitive advantage in the financial plan are an acquisition cost below EUR 1,137, one third of lifetime value; a cash burn that stays inside the seed through Year 2; and net revenue retention above 100% before the Series A, meaning existing customers spend more on upgrades than the firm loses to cancellations. Those three factors will determine the investor case, drive expansion revenue, and sustain runway to breakeven. The product offered is both scarce and relevant in Spain, Italy and Portugal, and will therefore help Compliance.AI to establish a competitive advantage.

8. IMPLEMENTATION PLAN

Build to the deadline. Compound after it.

One date organises everything: 2 December 2027, when the deferred high-risk enforcement deadline takes effect. Every in-scope SME must hold compliant documentation by that morning. The whole plan is built backwards from it.

The backward logic fixes the launch date. A customer needs months, not days, to classify their AI systems, generate the documentation, and file it. Launching in August 2027 leaves a customer sixteen months to reach compliance before the deadline closes. Launch any later and the platform arrives too late to be useful; launch earlier and it ships before the product is safe to trust. August 2027 is the date the deadline allows, and that date fixes everything before it. Table 8.1 maps seven workstreams across five phases from incorporation through Year 5 profitability. Darker shading marks the build or ramp phase for each workstream; lighter shading marks steady-state operations.

8.1 The Critical Path: Build First, Sell Second

A customer cannot buy a product that does not exist. The pre-launch period, July to December 2026, has one job: build the minimum viable platform, the risk classifier and the documentation generator, and earn nothing while doing it. This is why Year 1 in the financial model holds only five selling months. The cost is real: six months of cash burn before the first euro of revenue arrives. The reason is structural. Regulatory software that misclassifies on the first day destroys the trust the whole business depends on, so the product is built before it is sold, not alongside.

Two workstreams run together through pre-launch because each one feeds the other. Regulatory intelligence maps the Act's obligations into rules a machine can read. The platform turns those rules into a working classifier. Map first, build second. A classifier built

Table 8.1

Implementation roadmap.

Workstream	Pre-launch <i>Jul–Dec 2026</i>	Year 1 <i>2027</i>	Year 2 <i>2028</i>	Year 3 <i>2029</i>	Years 4–5 <i>2030–2031</i>
Milestones		Aug: launch 2 Dec: deadline		Mo. 33: EBITDA breakeven	seed recovered
Platform & product	MVP: risk classifier + doc generator	Launch Aug; conformity & monitoring modules	Annex III templates; API layer	Multi-tenancy; enterprise tier	Platform maturity; M&A optionality
Regulatory intelligence	Map Act obligations; horizon-scan	Monthly digest; advisory board input	High-risk deep dives	Ongoing monitoring	Steady state
Sales & marketing	Spain & Portugal partner pipeline	30 gross customers; Italy launch Q2	France & Benelux; MDF with law firms	Inbound & channel scale	Enterprise & Series A prep
Team build	Francisco full-time; incorporation & seed	Head of Compliance Q1; Italy Lead Q2; advisor live	+3 FTEs: engineer, Iberia AE, CS Manager (Ch. 5)	Engineering & CS scale-up	Scaled team
Infra & security	Cloud stack; SOC 2 roadmap; GDPR controls	ISO 27001 prep; pen-test cycle 1	SOC 2 Type II; EU data residency	Infra optimisation	Steady state
Partnerships	Identify Iberian law firm + accountancy partners	Referral agreements ES/PT/IT signed	FR/BE/NL channel; Big Four pilot	Channel expansion	Strategic alliances
Finance & legal	SL incorporation; €400k seed; ENISA	Monthly burn vs. runway; seed deployed	EBITDA breakeven watch (month 33)	Series A preparation	Profitability; investor exit options

Seven workstreams across five phases from inception to profitability. Darker cells mark the build or ramp phase for each workstream; lighter cells mark steady-state operations. Abbreviations: FTE = full-time equivalent; MDF = market development funds; API = application programming interface; SL = *Sociedad Limitada*; M&A = mergers and acquisitions.



before the obligations are mapped would encode guesses, and guesses are the one thing a compliance tool cannot ship.

The launch in August 2027 then opens the five-month selling window that produces the 30 gross customers of the Year 1 base case (Chapter 7). Two further modules, conformity workflow and post-market monitoring, ship in the same year, so that a customer who classifies in August can document, file, and begin monitoring before the December 2027 deadline.

8.2 The Team Arrives in the Order the Work Demands

Each Year 1 hire is timed to a dependency, not to a calendar. The founder carries product, sales and regulatory contact from incorporation. Two hires then land in sequence, each one clearing the work that follows it (Chapter 5).

The Head of AI Act Compliance joins in the first quarter, before the conformity module ships. Article 14 of the Act requires a qualified human to review high-risk classifications, so the platform cannot offer that module without the hire in place. The sequence is forced by the regulation, not chosen for convenience. The Italy Lead joins in the second quarter, timed to the Italy launch in the same quarter, and converts ten of the thirty Year 1 customers from Milan inbound. Three further hires, an engineer, an Iberia Account Executive and a Customer Success Manager, follow in Year 2 once revenue covers them.

8.3 Security Compounds into a Sales Asset

Security work starts in pre-launch and pays off two years later. The cloud stack, the GDPR controls and EU data residency are in place from the first day, because a compliance product that leaks data has no claim to make. The larger certification, SOC 2 Type II, arrives in Year 2. The delay is not slack. The standard requires an observation window of six to twelve months before the certificate can be issued, so Year 2 is the earliest it can complete.

That certificate is worth the wait. It opens the enterprise tier, where buyers demand proof



of security before they sign, and it answers the cyber risk treated in Chapter 9. Security spending in Year 1 is therefore not a cost centre. It is the foundation the Year 2 enterprise revenue stands on.

8.4 The Sequence Funds Itself

Every phase is sized so the cash from the round carries it to the next. The EUR 400,000 seed covers operating burn of EUR 343,784, capital expenditure of EUR 26,000, and a EUR 30,216 Year 2 cash cushion, with no second raise required before EBITDA break-even in month 33 from inception. The build phase is funded by the seed, the ramp by the seed plus early revenue, and the steady state by the business itself.

The roadmap reaches profitability on the capital it starts with. Build to the deadline. Then let the platform compound.

9. CRITICAL RISKS AND CONTINGENCY PLAN

Name the risk. Pre-commit the response.

Chapter 7 stress-tests revenue; this chapter stress-tests strategy. Ten risks could break the plan. Four would kill it. We name them, score them, and pre-commit the response.

9.1 Risk Framework

Every risk is scored on two questions: how likely is it, and how much damage does it do if it happens? Each answer is Low, Medium or High, and the combination places the risk in one of three bands: Green (monitor), Amber (active mitigation) or Red (material). Figure 9.1 plots the ten risks the firm tracks. Four land in the Red band and receive their own treatment below. The full register, with owners and trigger metrics, sits in Appendix A.1.

9.2 Material Risks


9.2.1 Demand: enforcement delay and SME awareness (R1, R6)

The Digital Omnibus is the EU package that provisionally moves the enforcement deadline for high-risk AI systems (Annex III) to a later fixed date (Marrache & Sieniewicz, 2026). A later deadline does not destroy demand. It postpones it. The real exposure is cash: the firm must survive the gap between launch and the moment buying becomes unavoidable.

Three mitigations are pre-committed. First, Spain enforces regardless of Brussels: AESIA, the Spanish AI supervision agency, published 16 compliance guides in December 2025 and runs the only national AI sandbox in the EU (AESIA, 2025; Cooper &

Figure 9.1

Risk heat map: ten risks scored by likelihood and impact.



Red cells flag material risks (R1–R4) treated in Section 9.2. Full register with owners and triggers in Appendix A.1.

Oberschelp de Meneses, 2025b). Spanish SME demand holds whatever the Omnibus decides. Second, the content engine in Chapter 4 turns an obligation firms know is coming into one they act on now. Already, 33% of EU tech SMEs are stripping AI features because they are unsure of the rules (Association for Competitive Technology, 2025), and those firms read regulatory guidance before they read pricing pages. Third, a trigger: formal adoption of the Omnibus with a confirmed long deferral moves the NIS2 module onto the Year 1 roadmap, so the firm also sells against a regulation already in force.

The deadline moves. The Article 99 fine cap of EUR 15 million does not (Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act), 2024).

9.2.2 Competition: generalist downmarket entry (R2)

Vanta or Drata, the two largest compliance-automation platforms, could ship an AI Act tier for SMEs inside 18 months. It has happened before: Drata closed Vanta's SOC 2 feature lead in 24 months (Sacra, 2024a, 2024b). The threat is real.

Two structural defences hold. The first is cost. A EUR 199 per month customer cannot repay an enterprise sales team: at contracts below USD 5,000 per year, the cost of winning a customer must be earned back within nine months (Benchmarkit, 2025b; Key-Banc Capital Markets, 2024), and enterprise cost structures cannot do that. The second is data. Every conformity assessment processed in Year 1 records how AESIA, ACN and ANACOM actually interpret the Act, knowledge a late entrant cannot buy or back-fill. Doubling cumulative case experience cuts error rates by 10–25% in regulated knowledge work (Argote, 2013).

Operational mitigations stack on top. Annual contracts lock in the Year 1 cohort before the competitive window opens. NIS2 and CRA modules ship in Year 2, so switching away means replacing three products, not one. Exclusivity agreements with two to three mid-tier EU law firms lock distribution. Trigger: any major GRC vendor announcing an AI Act SME tier pulls the Series A forward by six months to defend the position.

We get there first.

9.2.3 Technology and liability: classification error (R3)

Large language models peak at roughly 70% accuracy on legal reasoning (Fan et al., 2026; Fei et al., 2024). Three answers in ten can be wrong. A misclassification that lets a customer ship an AI system the Act actually covers is the worst Tuesday this firm can have: an Article 99 fine for the customer, an indemnity claim against the platform, and a reputation absorbing both inside a 700-firm reachable market.

The architecture absorbs the error by design. Article 14 of the EU AI Act requires human oversight of high-risk classification (Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act), 2024), so the human review layer is a legal requirement the firm builds in, not a cost it apologises for. The Head of AI Act Compliance hired in Year 1 (Chapter 5) reviews every borderline case. The terms of service draw the legal line: Compliance.AI provides the tool, while the customer remains the deployer responsible under Articles 26 and 27 for how the AI system is used. Professional indemnity (PI) insurance of EUR 2–5 million caps what a claim can cost. Trigger: any customer regulatory escalation unresolved beyond five working days, already a Year 1 performance target.

9.2.4 Key person: founder concentration (R4)

One technical founder carries product, sales and regulator contact until the Series A. That is the firm's single point of failure, and the plan treats it as one. The founder's zero-salary period ends at the earliest of EUR 240,000 in annual recurring revenue, any external financing event, or month 18 (Index Ventures, 2017). A fixed expiry date keeps the discipline honest.

Four mitigations are in place. Miguel de Reina's advisory engagement covers Southern European enterprise introductions at zero Year 1 cash cost. Written, asynchronous operations in Notion (Chapter 6) keep process on the page, not in the founder's head. Equity grants to the Head of Compliance (1.0%) and the Italy Lead (1.0%), vesting over four years, tie retention to the Series A horizon (Carta, 2024). The Year 2 option pool reserves

equity for a co-founder if bandwidth fails. Trigger: two consecutive customer-facing weeks lost, or quarterly objectives (OKRs) missed by more than 20% for two quarters running, starts the co-founder search.

The firm runs on three people in Year 1 (Lepak & Gowan, 2017). Two of them are paid.

9.3 Sensitivity Analysis

Customer acquisition is stress-tested in Appendix A.5, Table 9.7: Bull at EUR 197,104 ARR, Bear at EUR 49,276, Downside at EUR 26,281. The input not tested there is churn. The base model assumes 3.5% of customers cancel each month, inside the 3–7% benchmark range for SME software (High Alpha, 2024; UserMotion, 2024). At 5% monthly churn, Year 1 net customers fall from 24 to 19 and the lifetime value of an Essentials customer drops from EUR 3,411 to EUR 2,388. That breaches the 3:1 lifetime-to-acquisition floor once the fully loaded acquisition cost passes EUR 800. At 7%, the unit economics break entirely. Churn is the number the firm watches hardest from the first cohort onward.

9.4 Exit and Pivot Triggers

The exit and expansion triggers are quantified in Chapter 7, Section 7.5. A sustained 25% fall in monthly customer additions across two consecutive quarters, or a confirmed further delay of six months or more beyond 2 December 2027, activates the focused exit: a single founder on zero salary and the sale of the firm or its technology to a GRC platform seeking AI Act content. Growth at 1.5 to 2 times the base case activates an accelerated Series A.

Three pivot options exist if AI Act demand fails entirely. A NIS2 or CRA module sells a different regulation to the same SME segment. A move upmarket to the 250–500 employee band wins larger contracts but gives up the specialist position that keeps the incumbents out. A hybrid model of software plus advisory trades defensibility for revenue. The base case needs none of them. The plan should still know what they are.

Six load-bearing assumptions carry the plan, each with a failure mode and a validation path in Appendix A.2. The base case is built to survive its own bear.

REFERENCES

- AESIA. (2025, December 16). *Guidelines published to support compliance with the AI act*. Retrieved April 25, 2026, from <https://aesia.digital.gob.es/en/present/20251216-guidelines-published-to-support-compliance-with-the-ai-act>
- AI, A. (2026). *Pricing* [Claude API docs]. Retrieved April 22, 2026, from <https://platform.claude.com/docs/en/about-claude/pricing>
- Argote, L. (2013). *Organizational learning: Creating, retaining and transferring knowledge* (2nd ed.). Springer. <https://doi.org/10.1007/978-1-4614-5251-5>
- Association for Competitive Technology. (2025, October). *The hidden cost of AI regulations for EU and UK startups and SMEs*. Association for Competitive Technology. Retrieved April 15, 2026, from <https://actonline.org/the-hidden-cost-of-ai-regulations-a-survey-of-eu-uk-and-u-s-companies/>
- Benchmarkit. (2025a). *2025 SaaS performance metrics benchmark report*. Retrieved April 15, 2026, from <https://www.benchmarkit.ai/2025benchmarks>
- Benchmarkit. (2025b). *2025 SaaS performance metrics benchmarks*. Benchmarkit. Retrieved April 25, 2026, from <https://www.benchmarkit.ai/2025benchmarks>
- Bessemer Venture Partners. (2024). *State of the cloud 2024*. Retrieved April 15, 2026, from <https://www.bvp.com/atlas/state-of-the-cloud-2024>
- Bitkom e.V. (2025, September 10). *Breakthrough in artificial intelligence*. Retrieved June 16, 2026, from <https://www.privacy-conference.com/en/news/breakthrough-artificial-intelligence>
- Bonsignore, R., Scassellati-Sforzolini, G., Checcacci, G., Bosio, G. D., Mammì Borruto, F., & Onorato, P. M. (2025, October 30). *Italy adopts the first national AI law in europe complementing the EU AI act*. Retrieved May 6, 2026, from <https://www.clearyiptechinsights.com/2025/10/italy-adopts-the-first-national-ai-law-in-europe-complementing-the-eu-ai-act/>
- Carroll, G. R. (1985). Concentration and specialization: Dynamics of niche width in populations of organizations. *American Journal of Sociology*, 90(6), 1262–1283. <https://doi.org/10.1086/228210>

- Carroll, G. R., Dobrev, S. D., & Swaminathan, A. (2002). Organizational processes of resource partitioning. *Research in Organizational Behavior*, 24, 1–40. [https://doi.org/10.1016/S0191-3085\(02\)24002-2](https://doi.org/10.1016/S0191-3085(02)24002-2)
- Carta. (2024). *State of private markets q4 2024*. Retrieved April 15, 2026, from <https://carta.com/blog/state-of-private-markets-q4-2024/>
- ComplyJet. (2025). *Vanta pricing guide 2025: Real costs, ROI, and hidden fees*. Retrieved April 30, 2026, from <https://www.complyjet.com/blog/vanta-pricing-guide-2025>
- Cooper, D., & Oberschelp de Meneses, A. S. (2025a). *Inside privacy updates on developments in data privacy and cybersecurity*. Retrieved April 25, 2026, from <https://www.insideprivacy.com/artificial-intelligence/spain-issues-guidance-under-the-eu-ai-act/>
- Cooper, D., & Oberschelp de Meneses, A. S. (2025b, December 18). *Spain issues guidance under the EU AI act*. Retrieved April 19, 2026, from <https://www.insideprivacy.com/artificial-intelligence/spain-issues-guidance-under-the-eu-ai-act/>
- Council of the European Union. (2026, May 7). *Artificial intelligence: Council and parliament agree to simplify and streamline rules*. Retrieved June 9, 2026, from <https://www.consilium.europa.eu/en/press/press-releases/2026/05/07/artificial-intelligence-council-and-parliament-agree-to-simplify-and-streamline-rules/>
- Credo AI. (2025). *Credo AI named a leader in 2025 AI governance solutions report* [Businesswire]. Retrieved April 15, 2026, from <https://www.businesswire.com/news/home/20250825500906/en/Credo-AI-Named-a-Leader-in-2025-AI-Governance-Solutions-Report>
- DiMaggio, P. J., & Powell, W. W. (1983). The iron cage revisited: Institutional isomorphism and collective rationality in organizational fields. *American Sociological Review*, 48(2), 147–160. <https://doi.org/10.2307/2095101>
- Dyer, J. H., & Singh, H. (1998). The relational view: Cooperative strategy and sources of interorganizational competitive advantage. *The Academy of Management Review*, 23(4), 660–679. <https://doi.org/10.2307/259056>
- ENISA — Empresa Nacional de Innovación. (2025). *Línea de crédito digital: Financiación para empresas tecnológicas*. Retrieved April 20, 2026, from <https://www.enisa.es/es/financiacion/lineas-de-financiacion/linea-digital>

- EuroComply. (2026). *Platform and pricing overview* [EU compliance, distilled]. Retrieved April 15, 2026, from <https://eurocomply.app/>
- European Commission. (2025). *Standardisation of the AI act*. Retrieved April 25, 2026, from <https://digital-strategy.ec.europa.eu/en/policies/ai-act-standardisation>
- European DIGITAL SME Alliance. (2025). *AI act compliance made easier: Help is on its way for SMEs developing AI solutions*. Retrieved April 15, 2026, from <https://www.digitalsme.eu/ai-act-compliance-made-easier-help-is-on-its-way-for-smes-developing-ai-solutions/>
- European Digital SME Alliance. (2025). *EU SME digital performance benchmarks 2025*. European Digital SME Alliance. Retrieved June 9, 2026, from <https://www.digitalsme.eu/>
- Eurostat. (2025). *Use of artificial intelligence in enterprises* [Statistics explained]. Retrieved April 15, 2026, from https://ec.europa.eu/eurostat/statistics-explained/index.php?title=Use_of_artificial_intelligence_in_enterprises
- Fan, Y., Ni, J., Merane, J., Tian, Y., Hermstrüwer, Y., Huang, Y., Akhtar, M., Salimbeni, E., Geering, F., Dreyer, O., Brunner, D., Leippold, M., Sachan, M., Stremitzer, A., Engel, C., Ash, E., & Niklaus, J. (2026, April 2). LEXam: Benchmarking legal reasoning on 340 law exams. <https://doi.org/10.48550/arXiv.2505.12864>
- Fei, Z., Shen, X., Zhu, D., Zhou, F., Han, Z., Zhang, S., Chen, K., Shen, Z., & Ge, J. (2024). LawBench: Benchmarking legal knowledge of large language models. Retrieved April 25, 2026, from <https://arxiv.org/abs/2309.16289>
- Foundation Marketing. (2024). *How vanta built a \$ 128k/month SEO moat*. Retrieved April 30, 2026, from <https://foundationinc.co/lab/vanta>
- GDPR.eu. (2019). *Millions of small businesses aren't GDPR compliant: Our survey finds*. Retrieved May 5, 2026, from <https://gdpr.eu/2019-small-business-survey/>
- Glassdoor. (2024). *Compliance manager and AI governance officer salaries in europe*. Retrieved June 9, 2026, from <https://www.glassdoor.com/Salaries/>
- Grand View Research. (2025). *RegTech market size, share and trends: Industry report 2033*. Retrieved April 15, 2026, from <https://www.grandviewresearch.com/industry-analysis/regulatory-technology-market>
- Grant, R. M. (2010). *Contemporary strategy analysis: Text and cases* (7th ed.). Wiley-Blackwell.

- Hays. (2025). *Guía salarial hays 2025: Salarios en españa*. Retrieved April 15, 2026, from <https://www.hays.es/guia-salarial>
- Hernández-Perlines, F., Araya-Castillo, L., Palos-Sánchez, P., & Charm, H. B. (2024). Can hybrid strategy improve SME performance? the role of innovation and adaptive capacity. *European Journal of Innovation Management*, 27(9), 173–196. <https://doi.org/10.1108/EJIM-09-2023-0803>
- High Alpha. (2024). *SaaS benchmarks 2024*. Retrieved April 15, 2026, from <https://www.higalphalpha.com/saas-benchmarks/2024>
- Holistic AI. (2025). *Platform overview and EU AI act risk calculator*. Retrieved April 15, 2026, from <https://www.holisticai.com>
- Index Ventures. (2017). *Rewarding talent: A guide to stock options for european entrepreneurs* [tex.pubstate: Continually updated; latest content accessed 2026-05-06]. Retrieved May 6, 2026, from <https://www.indexventures.com/rewarding-talent/>
- International Association of Privacy Professionals. (2025). *AIGP certification: Artificial intelligence governance professional exam fees and study resources*. Retrieved May 6, 2026, from <https://iapp.org/certify/aigp/>
- KeyBanc Capital Markets. (2024). *2024 private SaaS company survey*. KeyBanc Capital Markets. Retrieved April 25, 2026, from <https://www.key.com/businesses-institutions/industry-expertise/2024-saas-survey-results.html>
- Kim, E., Nam, D.-i., & Stimpert, J. L. (2004). The applicability of porter's generic strategies in the digital age: Assumptions, conjectures, and suggestions. *Journal of Management*, 30(5), 695–726. <https://doi.org/10.1016/j.jm.2003.12.001>
- Kornutick, L. (2026, February 17). *Global AI regulations fuel billion-dollar market for AI governance platforms* [Gartner]. Retrieved April 15, 2026, from <https://www.gartner.com/en/newsroom/press-releases/2026-02-17-gartner-global-ai-regulations-fuel-billion-dollar-market-for-ai-governance-platforms>
- Laurer, M., Renda, A., & Yeung, T. (2021, September 24). *Clarifying the costs for the EU's AI act* [CEPS]. Retrieved April 15, 2026, from <https://www.ceps.eu/clarifying-the-costs-for-the-eus-ai-act/>
- Legge 23 settembre 2025, n. 132 – Disposizioni e deleghe al Governo in materia di intelligenza artificiale (2025, September 23). Retrieved May 6, 2026, from https://www.gazzettaufficiale.it/atto/serie_generale/caricaDettaglioAtto/originario?atto.

dataPubblicazioneGazzetta=2025-09-25&atto.codiceRedazionale=25G00143&elenco30giorni=true

- Lepak, D. P., & Gowan, M. (2017). *Human resource management: Managing employees for competitive advantage* (3rd ed.). Chicago Business Press.
- Lighter Capital. (2024). *SaaS gross margins and how to increase yours*. Retrieved April 15, 2026, from <https://www.lightercapital.com/blog/saas-gross-margins-and-how-to-increase-yours>
- Marrache, E., & Sieniewicz, S. (2026). *EU digital omnibus AI update: Council and parliament agreed positions ahead of trilogue*. Retrieved April 19, 2026, from <https://www.addleshawgoddard.com/en/insights/insights-briefings/2026/technology/eu-digital-omnibus-ai-update-council-parliament-agreed-positions/>
- OpenView Partners. (2022). *Product-led growth and freemium conversion benchmarks*. Retrieved April 15, 2026, from <https://www.lennysnewsletter.com/p/what-is-a-good-free-to-paid-conversion>
- Osano. (2025). *Law firm referral programme* [Osano]. Retrieved April 15, 2026, from <https://www.osano.com/company/partners-resellers/law-firms>
- Osterwalder, A., Pigneur, Y., Bernarda, G., & Smith, A. (2014, October). *Value proposition design: How to create products and services customers want*. Wiley. <https://www.wiley.com/en-us/Value+Proposition+Design-p-9781118968055>
- Pellerin, J. (2026, February 1). *Credo AI pricing and review 2026* [CO-AIMS]. Retrieved April 15, 2026, from <https://co-aims.com/blog/credo-ai-review-2026-compliance-officers>
- PitchBook. (2026). *Credo AI company profile* [Pitchbook]. Retrieved April 15, 2026, from <https://pitchbook.com/profiles/company/459035-74>
- Porter, M. E. (1980). *Competitive strategy: Techniques for analyzing industries and competitors*. Free Press.
- Porter, M. E. (1985). *Competitive advantage: Creating and sustaining superior performance*. Free Press.
- PricewaterhouseCoopers. (2025). *Spain: Individual — other taxes*. Retrieved April 10, 2026, from <https://taxsummaries.pwc.com/spain/individual/other-taxes>
- ProductLed. (2024). *Product-led growth benchmarks*. Retrieved April 15, 2026, from <https://productled.com/blog/product-led-growth-benchmarks>

- PwC Worldwide Tax Summaries. (2025a). *Italy: Individual – other taxes (social security contributions)*. Retrieved May 6, 2026, from <https://taxsummaries.pwc.com/italy/individual/other-taxes>
- PwC Worldwide Tax Summaries. (2025b). *Portugal: Individual – other taxes (social security contributions)*. Retrieved May 6, 2026, from <https://taxsummaries.pwc.com/portugal/individual/other-taxes>
- Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act) (2024, July 12). Retrieved April 25, 2026, from <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>
- Sacra. (2024a). *Drata: Continuous compliance for SMB SaaS*. Retrieved April 25, 2026, from <https://sacra.com/c/drata/>
- Sacra. (2024b). *Vanta: Compliance automation at scale*. Retrieved April 25, 2026, from <https://sacra.com/c/vanta/>
- Scott, W. R. (2001). *Institutions and organizations: Ideas, interests, and identities* (2nd ed.). Sage.
- Secure Privacy. (2026, March 23). *Cost of GDPR compliance: A realistic breakdown for 2026*. Retrieved May 5, 2026, from <https://secureprivacy.ai/blog/cost-of-gdpr-compliance>
- Startup Portugal. (2025). *Vouchers para startups – novos produtos digitais/tecnológicos*. Retrieved April 20, 2026, from <https://startupportugal.com/pt/programs/vouchers-para-startups-novos-produtos-digitais-tecnologicos/>
- UserMotion. (2024). *SaaS churn rate benchmark 2024*. Retrieved April 15, 2026, from <https://usermotion.com/saas-churn-rate-benchmark-2024>
- Vanta. (2025, July 23). *Vanta raises \$ 150m series d to power the future of AI-driven trust*. Retrieved June 16, 2026, from <https://www.businesswire.com/news/home/20250723901336/en/Vanta-Raises-150M-Series-D-to-Power-the-Future-of-AI-Driven-Trust>

APPENDIX

This appendix collects the supporting detail referenced throughout the plan. Sections A.1 and A.2 document the risk register and the assumptions audit behind Chapter 9. Sections A.3–A.5 hold the financial model behind Chapter 7: the model assumptions, the five-year statements, and the scenario grid. Each section is self-contained and cross-referenced from the chapter it supports.

A.1 Risk Register

Ten risks tracked by the firm. Likelihood and impact scored qualitatively (Low, Medium, High); band derived from the heat map in Figure 9.1. Material risks (Red band) are treated in Chapter 9, Section 9.2; the remainder sit under active monitoring with the owners and triggers listed below.

Mitigation cross-reference:

- R1, R6** Spanish enforcement holds; content engine converts latent obligation (Chapter 9, Section 9.2.1; Chapter 4).
- R2** Annual contracts; NIS2 / CRA bundle; partner exclusivity (Section 9.2.2).
- R3** Article 14 human review layer; PI insurance; AWS Frankfurt residency (Section 9.2.3; Chapter 6).
- R4** Miguel advisory; async ops; co-founder option pool (Section 9.2.4; Chapter 5).
- R5** Bundle pricing across classification, workflow, and Article 72 monitoring (Chapter 3).
- R7** Quarterly commission review; backup direct outreach by founder (Chapter 5).
- R8** SOC 2 Type II in Year 2; 14-day content cycle; EU residency by default (Chapter 6).
- R9** Bridge facility from Spanish public co-investment (ENISA, CDTI); reduced burn at single-founder mode (Chapter 7).

Table 9.1

Risk register: ten tracked risks.

ID	Risk	Category	Owner	L	I	Band	Trigger metric
R1	Enforcement timing uncertain pending Omnibus adoption	Strategic	CEO	M	H	Red	Digital Omnibus Official Journal publication
R2	Vanta / Drata ship AI Act SME tier	Strategic	CEO	M	H	Red	GRC vendor announcement
R3	LLM error causes customer audit failure	Operational	Head Compliance	M	H	Red	Escalation unresolved > 5 days
R4	Founder bandwidth collapse	Operational	Board	M	H	Red	2 weeks lost / 2 quarters OKR miss
R5	Article 11(1) template commoditises docs	Strategic	Head Compliance	L	M	Green	Commission publication
R6	SME deadline awareness stays low	Strategic	CEO	M	M	Amber	Lead conversion < 3% by M6
R7	Italy Lead misses Year 1 commission	Operational	CEO	M	M	Amber	Italian ARR < 40% target by Q3
R8	Cyber incident or data breach	Operational	CEO	L	H	Amber	Security incident severity 2+
R9	Seed round closes 6+ months late	Financial	CEO	L	H	Amber	Term sheets < 2 by Month 3 of raise
R10	Eurozone recession compresses SME IT	Macro	CEO	L	M	Green	EU SME ICT spend YoY decline

L = Likelihood; I = Impact. Material risks (R1–R4) receive deep treatment in Chapter 9, Section 9.2.

Amber risks under active mitigation are referenced through Sections 9.3–9.4. Green risk under monthly monitoring only.



R10 Pricing flexibility through monthly billing option; downside scenario reduces team to founder only (Chapter 7).

A.2 Key Assumptions Audit

The plan rests on six load-bearing assumptions. Each one carries a failure mode and a validation path. If two or more assumptions break simultaneously, the firm activates the exit triggers in Chapter 9, Section 9.4.

Table 9.2

Key assumptions, failure modes, and validation paths.

#	Assumption	Failure mode	Validation path
A1	54,000–66,000 AI-using SMEs; 7,000–15,000 Annex III in-scope in Spain, Italy, Portugal	In-scope universe halves; Year 3 base case still needs under 10% of it	Eurostat 2026 ICT survey (annual update)
A2	Founder closes 2 deals per month across 5 selling months in Year 1	Bear scenario activates: 15 gross customers; seed runway shortened by 6 months	First 90 days of pipeline conversion rate
A3	Customer-knowledge moat compounds inside an 18-month window	Generalist closes the feature gap before lock-in; NIS2 cross-sell accelerates to compensate	Monthly AESIA case-resolution rate vs. competitor content velocity
A4	Spain AESIA enforcement holds regardless of the Brussels deferral	Demand stalls until the deferred deadline nears; Bear scenario triggers exit strategy	Omnibus adoption; AESIA enforcement communications
A5	60% Year 1 gross margin with human-in-the-loop review at scale	Margins collapse to 40%; seed runway shrinks 4 months	Month 6 unit economics review on first 15 customers
A6	3.5% monthly churn ceiling holds for SME tier	LTV:CAC breaches 3:1; expansion revenue stalls; Series A valuation compresses	First cohort retention curve, month 6 onward

Sources: Chapter 2 (A1, A4); Chapter 7 (A2, A5); Chapter 3 (A3); Chapter 4 and Chapter 7 (A6).

Joint failure scenarios. The model tolerates single-assumption breaks. It does not tolerate joint breaks. Two combinations are explicitly modelled in Chapter 9, Section 9.4:

- **A2 + A4 break together** (Brussels delays enforcement AND founder underperforms acquisition target by 25%): activates Bear-to-Downside transition; exit strategy via acqui-hire.
- **A5 + A6 break together** (margins collapse AND churn exceeds 5%): unit economics fail at the Essentials tier; pricing reset to EUR 249 minimum plus mandatory annual contracts on all new customers.

A single assumption breaking activates the corresponding mitigation in the risk register (Appendix A.1). Two simultaneous breaks activate the exit triggers in Chapter 9, Section 9.4. Three or more is a wind-down.

A.3 Financial Model Assumptions

ARPU is average revenue per account per month. Net revenue retention is the share of a customer cohort's spending kept and expanded one year on. CAC is the all-in cost of acquiring one customer. Blended ARPU rises from EUR 339 in Year 1 to roughly EUR 445 by Year 5 as the tier mix shifts toward Professional and Enterprise and net revenue retention climbs from 95% to 112%; the customer counts in Appendix A.4 therefore pair with a rising ARPU, not the Year 1 figure. The Year 1 gross margin of 60% reflects human-in-the-loop documentation review while the Annex IV generator is calibrated in production; margin reaches 84% by Year 3 as multi-tenancy spreads fixed cost, then settles at 82% as enterprise accounts add bespoke review hours.

A.4 Five-Year Financial Statements

A.5 Scenario Detail

⁶Fully loaded Year 1 payroll of EUR 143,625 comprises EUR 116,325 fixed payroll (founder, Head of Compliance, Italy Lead), EUR 10,550 Italy Lead commission at the base-case target, EUR 5,000 advisory board cash, EUR 7,000 freelance legal review buffer, and EUR 4,750 hiring-cycle contingency.

**Table 9.3***Core financial model assumptions.*

Assumption	Value	Source / Confidence
Monthly logo churn (Y1)	3.5% declining to 1.8% (Y5)	UserMotion 2024 SMB cohort (UserMotion, 2024). MEDIUM
Gross margin (Y1 → Y5)	60% → 82%	Lighter Capital 2024; Benchmarkit 2025 (Benchmarkit, 2025a; Lighter Capital, 2024). HIGH
Net revenue retention (Y1 → Y5)	95% → 112%	Benchmarkit 2024 median 101% (Benchmarkit, 2025a). HIGH
Blended ARPU (Y1)	EUR 339/month	70/25/5 mix at EUR 199/499/1,499; Chapter 4 pricing. HIGH
Non-salary CAC (Y1)	EUR 400	EUR 12,000 S&M budget ÷ 30 gross customers. MEDIUM
Year 1 gross customers	30 (base); 60 (bull); 15 (bear); 8 (downside)	Founder sales capacity model; Appendix A.5. MEDIUM
Payroll (fully loaded, Y1)	EUR 143,625	Chapter 5 compensation table plus reconciliation footnote ⁶ . HIGH

Table 9.4
Income statement: base case (EUR).

	Y1	Y2	Y3	Y4	Y5
Net customers (eop)	24	117	301	546	843
ARR run rate	98,552	511,434	1,426,456	2,752,835	4,500,342
Revenue	49,276	308,772	989,963	2,134,787	3,708,519
Cost of Revenue	(19,700)	(78,000)	(161,000)	(358,000)	(655,000)
Gross Profit	29,576	230,772	828,963	1,776,787	3,053,519
Gross Margin	60%	75%	84%	83%	82%
Payroll (fully loaded)	(143,625)	(280,000)	(480,000)	(720,000)	(950,000)
Sales & Marketing	(12,000)	(55,000)	(140,000)	(280,000)	(420,000)
General & Admin	(55,000)	(85,000)	(120,000)	(170,000)	(230,000)
EBITDA	(181,049)	(189,228)	88,963	606,787	1,453,519
EBITDA Margin	(367%)	(61%)	9%	28%	39%
Depreciation & amort.	(2,500)	(7,500)	(8,500)	(9,500)	(15,000)
EBIT	(183,549)	(196,728)	80,463	597,287	1,438,519
Interest expense	(2,000)	(2,500)	(2,500)	(2,500)	(3,000)
Profit before tax	(185,549)	(199,228)	77,963	594,787	1,435,519
Tax	0	0	0	0	0
Net Income/(Loss)	(185,549)	(199,228)	77,963	594,787	1,435,519

Cost of revenue is infrastructure, LLM API, and human-in-the-loop documentation review. ARR run rates reflect a blended ARPU rising from EUR 339 (Year 1) to roughly EUR 445 (Year 5) through tier-mix shift and expansion revenue (Appendix A.3). The model carries nil cash tax through Year 3: the EUR 384,777 of Year 1–2 carry-forward losses fully shield the Year 3 profit and part of Year 4. Once those losses are exhausted in Year 4, Spain's 15% rate for new companies (PricewaterhouseCoopers, 2025) applies. The net income line above is stated pre-tax from Year 4; a post-tax view lowers Year 4 by about EUR 43,000 and Year 5 by about EUR 215,000 at the 15% rate, funded from operating cash flow.

Table 9.5

Cash flow statement: indirect method (EUR).

	Y1	Y2	Y3	Y4	Y5
Net income/(loss)	(185,549)	(199,228)	77,963	594,787	1,435,519
(+) Depreciation & amort.	2,500	7,500	8,500	9,500	15,000
(–) Incr. receivables	(6,075)	(31,993)	(83,982)	(141,143)	(194,022)
(+) Incr. deferred rev.	9,855	41,288	91,503	132,638	174,750
(+) Incr. payables	7,126	10,792	16,685	31,808	40,849
Cash from operations	(172,143)	(171,641)	110,669	627,590	1,472,096
Capital expenditure	(8,000)	(18,000)	(15,000)	(18,000)	(30,000)
Seed financing	400,000	0	0	0	0
Net change in cash	219,857	(189,641)	95,669	609,590	1,442,096
Cash, start of year	0	219,857	30,216	125,885	735,475
Cash, end of year	219,857	30,216	125,885	735,475	2,177,571

Working capital: receivables at 45 days of revenue; deferred revenue at 10% of the ARR run rate; payables at 30 days of non-payroll cost. The minimum cash balance is EUR 30,216 at the end of Year 2.

Table 9.6
Balance sheet: year-end (EUR).

	Y1	Y2	Y3	Y4	Y5
Assets					
Cash & equivalents	219,857	30,216	125,885	735,475	2,177,571
Accounts receivable	6,075	38,068	122,050	263,193	457,215
PP&E + capitalised dev, net	5,500	16,000	22,500	31,000	46,000
Total assets	231,432	84,284	270,435	1,029,668	2,680,786
Liabilities					
Long-term debt (ENISA)	50,000	50,000	50,000	50,000	50,000
Deferred revenue	9,855	51,143	142,646	275,284	450,034
Accounts payable & accrued	7,126	17,918	34,603	66,411	107,260
Total liabilities	66,981	119,061	227,249	391,695	607,294
Equity					
Paid-in capital	325,000	325,000	325,000	325,000	325,000
Capital grant reserve	25,000	25,000	25,000	25,000	25,000
Accumulated retained earnings	(185,549)	(384,777)	(306,814)	287,973	1,723,492
Total equity	164,451	(34,777)	43,186	637,973	2,073,492
Total liab. + equity	231,432	84,284	270,435	1,029,668	2,680,786

**Table 9.7***Four-scenario summary: Year 1 and Year 3 exit metrics (EUR).*

	Bull	Base	Bear	Downside
	60 cust.	30 cust.	15 cust.	8 cust.
Y1 gross new customers	60	30	15	8
Y1 net customers (eop)	48	24	12	6
Y1 ARR run rate	197,104	98,552	49,276	26,281
Y3 net customers (eop)	491	301	155	90
Y3 ARR run rate	2,329,160	1,426,456	733,588	428,854
Y3 EBITDA	1,048,000	88,963	(176,000)	(300,000)
Months to EBITDA breakeven	18	33	44	>60
Seed required	350,000	400,000	500,000	600,000

Bull trigger: enforcement approaching the deferred deadline drives a demand surge. Bear trigger: a further enforcement delay of six months or more stalls purchase decisions. Downside trigger: founder bandwidth collapses under the dual engineering and sales load. Seed figures for the Bear and Downside columns state the capital those cases would have required had they been the founding plan; under the base-case EUR 400,000 seed, a sustained shortfall of that size activates the exit triggers in Chapter 9 before either case fully develops. Sources: cohort model calibrated against High Alpha (2024) and UserMotion (2024).